


MODULE 37

PostgreSQL Database Fundamentals

Design normalized PostgreSQL schemas from first principles -- tables, types, keys, constraints, indexes, and real connectivity from psql to Java.





MODULE 37 OVERVIEW

Design robust, normalized PostgreSQL databases and connect them to real applications.

Northstar CRM needs a real relational schema behind every screen -- customers, accounts, and addresses, modeled and constrained in PostgreSQL from day one.

DURATION & LAB



1.5 Hours Theory

Relational fundamentals, data types, keys, constraints, relationships, indexes, sequences, and connectivity.



2 Hours Lab

Design and build a real PostgreSQL schema for the CRM customer and account platform.



Scenario

Northstar CRM: CUSTOMER, ACCOUNT, ADDRESS, and status history, modeled and constrained end to end.

MODULE ARC



Understand the Model

Relational concepts, PostgreSQL architecture, and how data is organized into databases and schemas.



Master Keys & Constraints

Primary keys, foreign keys, uniqueness, NOT NULL, CHECK, and normalized relationships.



Connect and Govern

Indexes, sequences, psql, JDBC, users, roles, transactions, and enterprise design practice.

KEY TAKEAWAY Total time: 3.5 hours (1.5 hours theory + 2 hours hands-on lab). By the end, you will design, constrain, and connect to a normalized PostgreSQL schema.

WHY RELATIONAL DATABASES MATTER

Learn how to design robust, normalized databases in PostgreSQL for real enterprise systems.



Design Better Databases

Model entities, attributes, and relationships before writing a single CREATE TABLE.



Ensure Data Integrity

Enforce rules and consistency so the database itself rejects invalid data.



Support Scalable Apps

Well-structured schemas hold up as data and traffic grow.



Make Smarter Decisions

Clean, trustworthy data powers better reporting and analytics.

KEY TAKEAWAY A well-designed PostgreSQL schema is the single foundation every reliable, secure, high-performance application is built on.



Why Relational Databases Matter

Relational databases are the backbone of modern applications and enterprise systems.



Data Integrity

Enforce rules, constraints, and relationships to keep data accurate and consistent.



ACID Transactions

Ensure reliable operations with Atomicity, Consistency, Isolation, and Durability.



Complex Relationships

Model real-world connections between data using foreign keys and relationships.



Powerful Querying

Use SQL to query, analyze, and report data efficiently.



Scalability

Handle growing data volumes and users with proven architectures and indexing.




Security

Role-based access, permissions, and auditing protect sensitive data.



Relational databases provide the structure, reliability, and performance needed to build secure, data-driven applications at scale.



MODULE LEARNING OBJECTIVES

By the end of this module, you will be able to design, constrain, and connect to a normalized PostgreSQL database.

1. Understand Fundamentals -- explain relational concepts and PostgreSQL's architecture.
2. Choose Data Types -- select the right PostgreSQL type for every column.
3. Create Structures -- create databases, schemas, and tables with confidence.
4. Enforce Integrity -- apply primary keys, foreign keys, and constraints correctly.
5. Model Relationships -- design one-to-one, one-to-many, and many-to-many links.
6. Normalize Data -- reduce redundancy while keeping queries practical.
7. Index for Performance -- apply B-Tree, unique, and composite indexes.
8. Connect Real Clients -- use psql, pgAdmin, and JDBC from a Java application.

KEY TAKEAWAY Goal: design a normalized, constrained, and correctly indexed PostgreSQL schema that a real Java application can connect to safely.



Learning Objectives

By the end of this module, you will be able to:



Understand PostgreSQL

Explain PostgreSQL architecture and its key components.



Design Database Structures

Create databases, schemas, and tables using appropriate data types and constraints.



Model Relationships

Implement one-to-one, one-to-many, and many-to-many relationships using foreign keys.



Apply Constraints and Keys

Use primary keys, foreign keys, unique, not null, check constraints, and default values to ensure data integrity.



Use Indexes and Sequences

Create and manage indexes, sequences, and identity columns for performance and data consistency.



Work with Users and Permissions

Manage database users, roles, and permissions using the principle of least privilege.



Apply Transactions and ACID Principles

Demonstrate how transactions and ACID properties ensure reliable and consistent data operations.



These fundamentals form the foundation for building robust, secure, and high-performance applications with PostgreSQL.



POSTGRESQL OVERVIEW

PostgreSQL is a powerful, open-source, standards-compliant object-relational database system.



Open Source

Free, transparent, community- and enterprise-driven development with no licensing cost.



Standards Compliant

Strong adherence to the SQL standard, plus rich PostgreSQL-specific extensions.



Extensible

Custom types, functions, operators, and extensions (e.g., PostGIS) plug directly into the engine.



ACID & MVCC

Full transactional integrity with Multi-Version Concurrency Control for high concurrency.

KEY TAKEAWAY PostgreSQL combines standards-compliant SQL, strong ACID guarantees, and rich extensibility -- the engine this course, and its labs, are built on.

PostgreSQL Overview



PostgreSQL is a **powerful, open source object-relational** database system. It is known for **reliability, data integrity, extensibility, and performance**.

Key Features



ACID Compliant

Ensures reliable transactions with full ACID support.



Advanced Data Types

Supports JSON/JSONB, arrays, UUID, full-text search, and more.



Extensible

Custom data types, functions, operators, and extensions.



High Performance

Powerful query planner, indexes, and concurrency control.



Reliability & Integrity

Strong data integrity, foreign keys, constraints, and crash recovery.



Open Source & Community

Large global community and continuous improvement.

Use Cases



Enterprise Applications



Analytics & Reporting



Web & Mobile Applications




AI/ML and Data Science



JSON/Document Workloads

PostgreSQL Architecture



Why Organizations Choose PostgreSQL



Proven Reliability



Enterprise Ready



Scalability



Extensibility



Cost Effective



PostgreSQL combines the **robustness** of enterprise databases with the **flexibility** of open source.

POSTGRESQL ARCHITECTURE

PostgreSQL uses a process-per-connection model with a shared storage engine.

1

Clients / Apps

Connect to PostgreSQL over the network using a driver such as JDBC.

2

Postmaster

The supervisor process authenticates each connection and spawns a backend.

3

Backend Process

One dedicated OS process per connection handles that session's queries.

4

Shared Memory & Storage


Shared buffers, the WAL, and a lock table are shared across all backend processes.

KEY TAKEAWAY The Postmaster spawns one dedicated backend process per client connection; all processes share buffers, the write-ahead log, and a lock table.

PostgreSQL Architecture



PostgreSQL uses a client-server architecture with a multi-process model for performance, reliability, and scalability.



Key Takeaways

- ✓ PostgreSQL uses a multi-process architecture for stability and performance.
- ✓ Each client connection has its own backend process.
- ✓ Shared memory and background processes work together to handle concurrent workloads efficiently.

CLIENT-SERVER DATABASE ARCHITECTURE

PostgreSQL always runs as a server process that clients connect to over the network -- never embedded in the application.


Role	Responsibility
Server (postgres)	Runs continuously, owns the data files, enforces every constraint and permission.
Client	psql, pgAdmin, DBeaver, or a Java app -- sends SQL, receives result sets.
Protocol	The PostgreSQL wire protocol carries queries and results over TCP, typically port 5432.
Network Boundary	Client and server can run on the same machine or across a network -- the protocol is identical either way.

KEY TAKEAWAY Every client -- human or application -- talks to the same PostgreSQL server process over the same wire protocol; the server is the single source of truth.

Client-Server Database Architecture



PostgreSQL uses a **client-server model** where client applications request data and the PostgreSQL server processes the request, manages the data, and returns the results.



KEY TAKEAWAYS



Clients communicate with the PostgreSQL server over the network using standard protocols.



The server handles query parsing, optimization, execution, and transaction management.



Data is stored on disk with strong durability, consistency, and reliability.



This architecture enables scalability, security, centralized management, and data integrity.

DATABASE VS SCHEMA

A PostgreSQL server can host many databases; each database can hold many schemas, each a namespace for tables.

Level	What It Is	Typical Use
Server / Cluster	One running PostgreSQL instance.	One installation, many databases.
Database	An isolated collection of schemas, tables, and data.	One per application, e.g. crmdb.
Schema	A named namespace of tables inside one database.	Group by module or tenant, e.g. public, billing.
Table	The actual rows and columns of data.	customer, account, address.

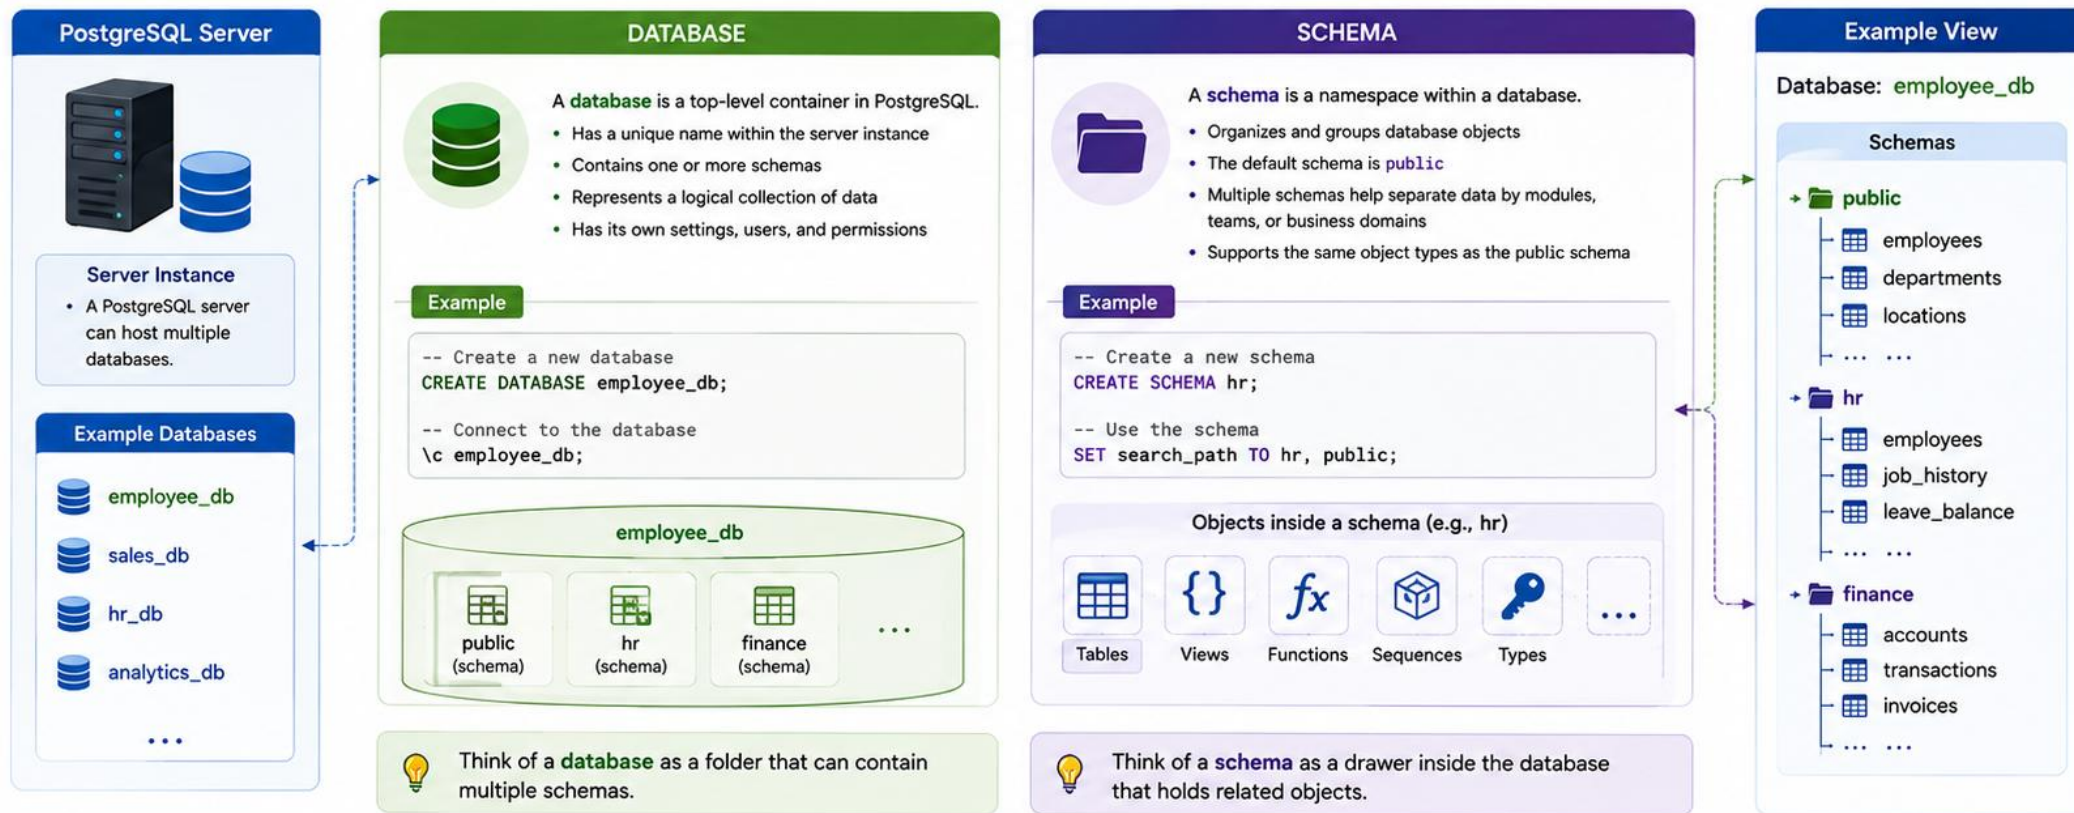
KEY TAKEAWAY A database isolates data completely; a schema inside it is just an organizational namespace -- most PostgreSQL apps use one database and the default public schema.

Database vs Schema



In PostgreSQL, a **database** is a container for schemas.

A **schema** is a namespace that contains database objects such as tables, views, and functions.



KEY TAKEAWAYS



A **database** is a container that can hold multiple schemas.



A **schema** is a namespace that organizes database objects.



Use multiple schemas to logically separate data by team or domain.



Both databases and schemas have their own permissions and security.

TABLES AND ROWS

Data in PostgreSQL is organized into tables; each table is made of rows, and each row is one record.

customer_id	full_name	status	created_at
1	Amina Khan	ACTIVE	2026-01-10
2	Ravi Singh	PROSPECT	2026-01-11

Table -- stores data about one entity or business object, e.g. customer.

Row (tuple) -- a single record or instance of that entity, e.g. one customer.

Column (attribute) -- a specific property of the entity, e.g. full_name or status.

Every row in a table shares the exact same set of columns.

KEY TAKEAWAY A table stores one entity; each row is one instance of it; every row shares the exact same columns -- the foundation for everything that follows this module.

Tables and Rows



In PostgreSQL, data is stored in **tables**. A table is made up of **rows** and **columns**. Each row represents a record, and each column represents a **field** (attribute).

KEY CONCEPTS



Table

A collection of related data organized in columns and rows.



Row (Record)

A single entry in a table. Each row has a unique set of values.



Column (Field)

A named attribute of the table. All values in a column are of the same data type.



Analogy

Think of a table as a spreadsheet: columns are the headers, rows are the data entries.

EXAMPLE: employees TABLE

PRIMARY KEY



COLUMNS (FIELDS)

employee_id (integer)	first_name (text)	last_name (text)	email (text)	department_id (integer)	hire_date (date)	salary (numeric)
1	John	Doe	john.doe@acme.com	10	2023-01-15	75000.00
2	Jane	Smith	jane.smith@acme.com	20	2023-02-20	68000.00
3	Michael	Brown	michael.brown@acme.com	10	2023-03-05	72000.00
4	Emily	Davis	emily.davis@acme.com	30	2023-03-18	65000.00
5	Daniel	Wilson	daniel.wilson@acme.com	20	2023-04-01	70000.00

ROWS
(RECORDS)

SQL EXAMPLE

```
CREATE TABLE employees (  
  employee_id SERIAL PRIMARY KEY,  
  first_name TEXT NOT NULL,  
  last_name TEXT NOT NULL,  
  email TEXT UNIQUE,  
  department_id INTEGER REFERENCES departments(department_id),  
  hire_date DATE NOT NULL,  
  salary NUMERIC(12,2)  
);
```

-- Each column has a name and data type
-- Each row stores values for these columns

WHY IT MATTERS



Structured Data

Organizes data in a structured and efficient way.



Easy Retrieval

Makes querying and filtering data simple and powerful.



Data Integrity

Constraints (like primary key, foreign key) help maintain accuracy.



Scalability

Designed to handle large volumes of data reliably.



KEY TAKEAWAYS



A table stores data in a structured format of rows and columns.



Each row is a record. Each column is a field with a specific data type.



Primary keys uniquely identify each row in the table.



This structure enables efficient data management and powerful querying in PostgreSQL.

COLUMNS AND DATA TYPES

Every column has a declared data type that constrains exactly what values it can hold.



Type Enforcement

PostgreSQL rejects a value that does not match its column's declared type at insert time.



Storage & Precision

The type determines how much space a value uses and how precisely it is stored.



Operator Behavior

A column's type decides which operators and functions are valid against it, e.g. date arithmetic.



Constraints Ride Along

NOT NULL, DEFAULT, and CHECK are attached per column alongside its type.

KEY TAKEAWAY The declared type of a column is a contract: PostgreSQL enforces it on every insert and update, and it drives storage, precision, and which operations are valid.

Columns and Data Types



In PostgreSQL, a table is made up of **columns** (fields).
Each column has a name and a **data type** that defines the kind of values it can store.

EXAMPLE TABLE: employees

COLUMNS (FIELDS)						
employee_id (INTEGER)	first_name (TEXT)	last_name (TEXT)	email (TEXT)	department_id (INTEGER)	hire_date (DATE)	salary (NUMERIC)
1	John	Doe	john.doe@acme.com	10	2023-01-15	75000.00
2	Jane	Smith	jane.smith@acme.com	20	2023-02-20	68000.00
3	Michael	Brown	michael.brown@acme.com	10	2023-03-05	72000.00
4	Emily	Davis	emily.davis@acme.com	30	2023-03-18	65000.00
5	Daniel	Wilson	daniel.wilson@acme.com	20	2023-04-01	70000.00

ROWS
(RECORDS)



Think of a table as a spreadsheet:
Columns are the headers, rows are the data entries.

KEY POINTS

- ✓ Each column has a name and a data type.
- ✓ The data type determines what kind of values the column can store.
- ✓ Choosing the right data type improves data integrity and performance.

COMMON POSTGRESQL DATA TYPES

CATEGORY	DATA TYPE	DESCRIPTION	EXAMPLE
Numeric	INTEGER	Whole numbers (4 bytes)	42
	BIGINT	Large whole numbers (8 bytes)	9223372036854775807
	NUMERIC(p,s)	Exact numeric with precision and scale	12345.67
	REAL / DOUBLE PRECISION	Floating-point numbers	3.14159
Character	TEXT	Variable-length character string	'Hello PostgreSQL'
	VARCHAR(n)	Variable-length with limit	'Alice'
	CHAR(n)	Fixed-length character string	'USA '
Boolean	BOOLEAN	Stores true or false	true / false
Date / Time	DATE	Calendar date (YYYY-MM-DD)	2024-05-20
	TIME	Time of day (HH:MM:SS)	14:30:00
	TIMESTAMP	Date and time without time zone	2024-05-20 14:30:00
	TIMESTAMPTZ	Date and time with time zone	2024-05-20 14:30:00+00
Other	UUID	Universally unique identifier	550e8400-e29b-41d4-a716-446655440000
	JSON / JSONB	JSON data (JSONB is binary JSON)	{ "name": "Alice" }
	BYTEA	Binary data	\xDEADBEEF



KEY TAKEAWAYS



Columns define the structure
of a table.



Data types ensure only valid
values are stored.



Strong data types improve
accuracy and reliability.



Proper data types lead to
better performance.

POSTGRESQL DATA TYPES OVERVIEW

PostgreSQL ships a rich built-in type system -- choosing the right one is a real design decision.

Category	Examples	Typical Use
Numeric	INTEGER, BIGINT, NUMERIC, REAL	Counts, IDs, exact money amounts.
Character	VARCHAR(n), TEXT, CHAR(n)	Names, emails, free text, fixed codes.
Boolean	BOOLEAN	True/false flags, e.g. is_active.
Date/Time	DATE, TIME, TIMESTAMPTZ, INTERVAL	Timestamps, durations, schedules.
Identifier	UUID	Globally unique, non-sequential public identifiers.
Semi-Structured	JSON, JSONB	Flexible, schema-light document data.

KEY TAKEAWAY Six broad type categories cover almost every column PostgreSQL schemas need -- the next six slides walk each one with real examples.



PostgreSQL Data Types Overview

PostgreSQL provides a rich set of native data types to store different kinds of data efficiently and safely.

123 Numeric Types

For numbers, integers, decimals, and arbitrary precision values.

Type	Example
SMALLINT	32767
INTEGER	2147483647
BIGINT	9223372036854775807
NUMERIC(p,s)	12345.6789
DECIMAL(p,s)	99.99
REAL	3.14159
DOUBLE PRECISION	2.718281828459045

A Character Types

For storing text, strings, and fixed-length character data.

Type	Example
CHAR(n)	'ABC'
VARCHAR(n)	'Hello World'
TEXT	'This is a text value of any length'
CITEXT	'Case Insensitive Text'

📅 Date and Time Types

For storing date, time, timestamp, and interval data.

Type	Example
DATE	2025-05-18
TIME	14:30:00
TIME WITH TIME ZONE	14:30:00+05:30
TIMESTAMP	2025-05-18 14:30:00
TIMESTAMP WITH TIME ZONE	2025-05-18 14:30:00+05:30
INTERVAL	'2 days 3 hours'

☑ Boolean Type

For storing true/false values.

Type	Example
BOOLEAN	true / false

{ } UUID Type

For storing universally unique identifiers.

Type	Example
UUID	550e8400-e29b-41d4-a716-446655440000

JSON JSON Types

For storing semi-structured and flexible JSON data.

Type	Example
JSON	'{"name": "John", "age": 30}'
JSONB	'{"role": "admin", "active": true}'

💡 Key Takeaways

- ✓ Choose the right data type to ensure data integrity and performance.
- ✓ Use NUMERIC/DECIMAL for exact financial or precise calculations.
- ✓ Use TEXT for large text data.
- ✓ Use TIMESTAMP WITH TIME ZONE for timezone-aware applications.
- ✓ Use UUID for globally unique identifiers.
- ✓ Use JSONB for efficient querying and indexing of JSON data.

★ Pro Tip

Use the most appropriate data type from the beginning—changing data types later can be expensive!



PostgreSQL is extensible. You can also create custom data types, domains, and enumerated types to fit your business needs.

NUMERIC TYPES

PostgreSQL offers both fast approximate types and exact decimal types -- picking the wrong one silently corrupts money.

Type	Range / Precision	Use For
SMALLINT	-32,768 to 32,767	Small counters, flags.
INTEGER	About -2.1B to 2.1B	Most whole-number columns, foreign keys.
BIGINT	About -9.2 quintillion to 9.2 quintillion	High-volume IDs, large counts.
NUMERIC(p,s)	Exact, user-defined precision/scale	Money, e.g. NUMERIC(19,2).
REAL / DOUBLE PRECISION	Approximate, floating-point	Scientific data -- never money.

KEY TAKEAWAY Use INTEGER or BIGINT for whole numbers and NUMERIC(p,s) for money -- never REAL or DOUBLE PRECISION, which are approximate and will silently round currency.



PostgreSQL Numeric Types

PostgreSQL provides several numeric data types to store exact numbers, approximate numbers, and high-precision decimal values.

Data Type	Aliases	Description	Storage / Size	Precision	Scale	Example Values
SMALLINT	int2	Small-range integer	2 bytes	Up to 5 decimal digits	0	-32768, 0, 32767
INTEGER	int4	Standard integer	4 bytes	Up to 10 decimal digits	0	-2,147,483,648 0, 2,147,483,647
BIGINT	int8	Large-range integer	8 bytes	Up to 19 decimal digits	0	-9,223,372,036,854,775,808 0, 9,223,372,036,854,775,807
NUMERIC(p,s)	decimal	Exact numeric with user-defined precision and scale	Variable (up to 131,072 digits)	1 – 131,072	0 – p	NUMERIC(10,2) = 12345.67 NUMERIC(20,6) = 123.456789
DECIMAL(p,s)	decimal	Alias for NUMERIC(p,s)	Variable (up to 131,072 digits)	1 – 131,072	0 – p	DECIMAL(10,2) = 12345.67
REAL	float4	Single-precision floating-point	4 bytes	≈ 6 – 7 decimal digits	N/A	3.14159, -2.5, 1.0e6
DOUBLE PRECISION	float8	Double-precision floating-point	8 bytes	≈ 15 – 16 decimal digits	N/A	3.141592653589793, -1.234567890123456e10
SERIAL	int4 (auto-increment)	Auto-incrementing integer (sequence-based)	4 bytes	Up to 10 decimal digits	0	1, 2, 3, 4, ...
BIGSERIAL	int8 (auto-increment)	Auto-incrementing big integer	8 bytes	Up to 19 decimal digits	0	1, 2, 3, 4, ...



Note: Use **INTEGER** / **BIGINT** for whole numbers, **NUMERIC** / **DECIMAL** for exact decimal values, and **REAL** / **DOUBLE PRECISION** for approximate numeric values.



Numeric Type Selection Guide

- Use **SMALLINT** for small ranges to save space.
- Use **INTEGER** for general-purpose whole numbers.
- Use **BIGINT** for very large whole numbers.
- Use **NUMERIC/DECIMAL** for financial, accounting, and high-precision data.
- Use **REAL** or **DOUBLE PRECISION** for scientific calculations where slight rounding is acceptable.



Best Practices

- ✓ Use the smallest appropriate type to optimize storage and performance.
- ✓ Use **NUMERIC** with appropriate precision and scale for monetary values.
- ✓ Avoid floating-point types for financial calculations.
- ✓ Avoid over-allocating precision and scale.
- ✓ Use **BIGSERIAL** for large, auto-incrementing keys.



Example Use Cases

- 📄 **SMALLINT:** Status codes, flags
- 👤 **INTEGER:** Age, quantity, count
- 💰 **NUMERIC(10,2):** Price, salary, invoice amount
- 🔬 **DOUBLE PRECISION:** Scientific measurements
- 🔑 **BIGSERIAL:** Primary key for large tables

CHARACTER TYPES

PostgreSQL's three character types behave almost identically -- the real-world default is VARCHAR or TEXT.

Type	Behavior	Recommendation
VARCHAR(n)	Variable length, up to n characters; rejects longer input.	Use when a meaningful maximum length exists, e.g. email VARCHAR(255).
TEXT	Variable length, unlimited.	Use for genuinely open-ended text, e.g. notes.
CHAR(n)	Fixed length; PostgreSQL pads shorter values with spaces.	Rarely needed -- reserve for fixed-width legacy codes.

KEY TAKEAWAY Prefer VARCHAR(n) when a real maximum length exists and TEXT when it doesn't -- PostgreSQL stores both identically, so there is no performance reason to prefer CHAR(n).



PostgreSQL Character Types

Character types store textual data such as letters, numbers, symbols, and strings. PostgreSQL provides both fixed-length and variable-length character data types.

Data Type	Aliases	Description	Storage / Size	Example Value	Notes
CHAR(n) (character)	character(n)	Fixed-length character string.	Exactly n characters (padded with spaces if shorter)	'ABC ' CHAR(5)	Use for values with constant length requirements.
VARCHAR(n) (character varying)	character varying(n), varchar(n)	Variable-length character string with maximum length n.	Up to n characters (only uses space needed)	'Hello World' VARCHAR(20)	Most common choice for general text data.
TEXT	text	Variable-length character string with no length limit.	Up to 1 GB of data	'This is a text value of any length'	Use for large text documents or unlimited strings.
CITEXT	citext (extension)	Case-insensitive text type.	Up to 1 GB of data	'Case Insensitive Text' citext	Case-insensitive comparisons and uniqueness.
BPCHAR(n) (blank padded)	bpchar(n)	Internal type used to store CHAR(n).	Exactly n characters (padded with spaces)	'ABC ' BPCHAR(5)	System type mapped from CHAR(n).
NAME	name	Fixed-length string used for object names (identifiers).	Up to 63 bytes	'customer_table' NAME	Used internally for database object identifiers.

✓ Key Takeaways

- ✓ Use **CHAR(n)** for fixed-length values.
- ✓ Use **VARCHAR(n)** for variable-length values with a defined limit.
- ✓ Use **TEXT** for large or unlimited-length strings.
- ✓ Use **CITEXT** for case-insensitive text comparisons.
- ✓ Names of database objects are stored using **NAME** type.



Length Semantics

CHAR(n)

Always stores exactly n characters.

CHAR(5) = 'AB'
is stored as 'AB '
(3 trailing spaces)

VARCHAR(n)

Stores only the actual characters used.

VARCHAR(5) = 'AB'
is stored as 'AB'
(no extra spaces)



Examples in Use

```
CREATE TABLE employee (  
  emp_code   CHAR(6),      -- Fixed-length employee code  
  emp_name   VARCHAR(100), -- Variable-length name  
  emp_bio    TEXT,         -- Long biography  
  email      CITEXT        -- Case-insensitive email  
);
```

```
SELECT 'abc' = 'ABC';           -- returns false  
SELECT 'abc'::citext = 'ABC'::citext; -- returns true
```

Note: CITEXT requires the extension:
`CREATE EXTENSION IF NOT EXISTS citext;`



Best Practices

- Prefer **VARCHAR(n)** over **CHAR(n)** unless length is fixed.
- Use appropriate length n to balance storage and validation.
- Use **TEXT** for long descriptions, notes, or documents.
- Use **CITEXT** when user input should be case-insensitive.

BOOLEAN TYPES

PostgreSQL's BOOLEAN type has three possible states: true, false, and NULL (unknown).



TRUE

Accepts true, 't', 'yes', 'y', '1', and more as literal input.



FALSE

Accepts false, 'f', 'no', 'n', '0', and more as literal input.



NULL (Unknown)

A boolean column can be NULL unless declared NOT NULL -- 'unknown', not false.



Use It Directly

Prefer BOOLEAN over an INTEGER or CHAR(1) flag column.

KEY TAKEAWAY A PostgreSQL BOOLEAN is three-valued -- true, false, or NULL -- so `is_active IS NULL` is a real, different condition from `is_active = false`.



PostgreSQL Boolean Types

Boolean types store true/false values and are commonly used for flags, status indicators, and conditional logic in PostgreSQL.



Boolean Data Type

BOOLEAN

The BOOLEAN type stores logical values:
true or **false**

Internal Storage	1 byte
Allowed Values	true, false
Default	false
Category	Boolean / Logical
Standard	SQL Standard compliant



Accepted Input Values

Evaluated As TRUE	Evaluated As FALSE
true	false
t	f
yes	no
y	n
on	off
1	0



Input is not case-sensitive. Extra spaces are ignored.



Key Takeaways

- ✓ BOOLEAN stores **true** or **false** values.
- ✓ Commonly used for flags, switches, and status indicators.
- ✓ Use BOOLEAN for clear and efficient logical conditions.
- ✓ Default value is **false** if not specified.
- ✓ Boolean expressions in queries return **true** or **false**.



Examples in Use

```
CREATE TABLE employee (  
  id          SERIAL PRIMARY KEY,  
  name        VARCHAR(100) NOT NULL,  
  is_active   BOOLEAN NOT NULL DEFAULT true,  
  email_verified BOOLEAN NOT NULL DEFAULT false  
);  
  
INSERT INTO employee (name, is_active, email_verified) VALUES  
  ('Alice', true, true),  
  ('Bob', false, false);
```



Query Examples

```
-- Select active employees  
SELECT * FROM employee WHERE is_active = true;  
  
-- Select inactive employees  
SELECT * FROM employee WHERE is_active = false;  
  
-- Count verified emails  
SELECT COUNT(*) FROM employee WHERE email_verified = true;  
  
-- Boolean expression in SELECT  
SELECT name, is_active, email_verified,  
       (is_active AND email_verified) AS can_login  
FROM employee;
```



Best Practices

- ▶ Use BOOLEAN for logical values (true/false).
- ▶ Avoid storing boolean values as text or integers.
- ▶ Name boolean columns with prefixes like **is_**, **has_**, **can_**, **should_**, etc.
- ▶ Set appropriate default values for flags.
- ▶ Use boolean columns in WHERE, CHECK, and CASE expressions.



Did You Know? BOOLEAN is part of the SQL standard and evaluated as true (1) or false (0) in expressions and conditions.

DATE AND TIME TYPES

PostgreSQL's date/time types range from a plain calendar date to a fully timezone-aware timestamp.

Type	Stores	Recommendation
DATE	Calendar date only, no time.	Birthdays, due dates with no time component.
TIME	Time of day only, no date.	Rarely used alone -- daily schedules.
TIMESTAMP	Date and time, no timezone.	Avoid for anything crossing timezones.
TIMESTAMPTZ	Date and time, stored in UTC, timezone-aware.	Default choice for created_at, updated_at, and event times.
INTERVAL	A span of time, e.g. '3 days'.	Durations and date arithmetic results.

KEY TAKEAWAY Default to TIMESTAMPTZ for any column that records when something happened -- it stores every value normalized to UTC and converts correctly for any client's timezone.



PostgreSQL Date and Time Types

PostgreSQL provides rich date and time types to store calendar dates, times of day, timestamps, and time intervals with high precision and time zone support.

Core Date and Time Types

Data Type	Description	Storage	Precision	Example Value	Notes
DATE	Calendar date (year, month, day)	4 bytes	1 day	2025-05-18	No time of day
TIME [(p)]	Time of day (without date)	8 bytes	1 microsecond (default p = 6)	14:30:00.123456	Range: 00:00:00 to 24:00:00
TIME WITH TIME ZONE (TIMETZ)	Time of day with time zone offset	12 bytes	1 microsecond (default p = 6)	14:30:00.123456+05:30	Stores time and time zone offset
TIMESTAMP [(p)] (without time zone)	Date and time (without time zone)	8 bytes	1 microsecond (default p = 6)	2025-05-18 14:30:00.123456	Does not store time zone
TIMESTAMP WITH TIME ZONE (TIMESTAMPTZ)	Date and time with time zone	8 bytes	1 microsecond (default p = 6)	2025-05-18 14:30:00.123456+05:30	Stored in UTC, converted on retrieval
INTERVAL	Time interval (duration)	16 bytes	1 microsecond	2 days 3 hours 15 minutes 10.123456 seconds	Can be positive or negative

Example Usage

```
-- Insert examples
INSERT INTO event (event_date, start_time, created_at, duration)
VALUES
  ('2025-05-18', '14:30:00', '2025-05-18 14:30:00', INTERVAL '2 hours 15 minutes');

-- Current date and time examples
SELECT CURRENT_DATE;           -- 2025-05-18
SELECT CURRENT_TIME;           -- 14:30:00.123456
SELECT CURRENT_TIMESTAMP;      -- 2025-05-18 14:30:00.123456
SELECT CURRENT_TIMESTAMP AT TIME ZONE 'UTC'; -- 2025-05-18 14:30:00.123456+00
```

Common Operations

Add Interval	timestamp + INTERVAL '1 day'
Subtract Interval	timestamp - INTERVAL '2 hours'
Date Difference	DATE '2025-05-18' - DATE '2025-05-10' -- returns 8 days
Extract Part	EXTRACT(YEAR FROM timestamp) EXTRACT(MONTH FROM date)
Truncate	DATE_TRUNC('day', timestamp)

Key Takeaways

- ✓ Use DATE for storing calendar dates.
- ✓ Use TIME for storing time of day.
- ✓ Use TIMESTAMP for date and time without time zone.
- ✓ Use TIMESTAMPTZ for date and time with time zone.
- ✓ PostgreSQL stores TIMESTAMPTZ in UTC internally.
- ✓ Use INTERVAL for durations and differences.

Time Zone Behavior



TIMESTAMP: no time zone conversion is applied.
TIMESTAMPTZ: converted to/from session time zone.

Best Practices

- Prefer TIMESTAMPTZ for any data that represents specific points in time.
- Store all timestamps in UTC; convert in the application layer.
- Be explicit about time zones in your applications.
- Use appropriate precision (p) only when needed.
- Avoid mixing TIMESTAMP and TIMESTAMPTZ in the same column for the same data.

PostgreSQL follows ISO 8601 standard for date and time formats.

Default timestamp precision is microsecond (6 digits).

Valid date range: 4713 BC to 5874897 AD.

UUID

A UUID is a 128-bit universally unique identifier -- ideal for public-facing IDs that must not reveal sequence or volume.



Globally Unique

Effectively collision-free across tables, databases, and even separate systems.



Non-Sequential

Does not reveal how many rows exist or the order they were created in.



Generated Two Ways

`gen_random_uuid()` in PostgreSQL, or generated in the application before insert.



Trade-Off

Larger (16 bytes) than BIGINT and not naturally sortable by creation time.

KEY TAKEAWAY Use UUID for identifiers exposed outside the database -- API responses, public URLs -- where an attacker or competitor should not be able to infer sequence or volume from the value.



PostgreSQL UUID

UUID (Universally Unique Identifier) is a 128-bit identifier used to uniquely identify records across systems, databases, and distributed environments.

What is UUID?

UUID is a 128-bit (16-byte) identifier that is designed to be globally unique.

- Extremely low probability of duplication
- Ideal for distributed systems
- Does not reveal information (non-sequential)
- Standard format defined by RFC 4122

Example UUID

550e8400-e29b-41d4-a716-446655440000

8 hex
digits

4 hex
digits

4 hex
digits

4 hex
digits

12 hex
digits

UUID in PostgreSQL

Data Type	uuid
Storage Size	16 bytes (128 bits)
Format	xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx
Case	Hex digits (0-9, a-f)
Hyphens	Displayed with hyphens (standard format)
Indexing	Fully indexable and efficient

Create Table Example

```
CREATE TABLE employee (  
  id uuid PRIMARY KEY DEFAULT gen_random_uuid(),  
  name text NOT NULL,  
  email text UNIQUE,  
  created_at timestamp DEFAULT now()  
);
```

Key Takeaways

- ✓ UUID provides globally unique identifiers.
- ✓ Use `gen_random_uuid()` for random UUIDs in PostgreSQL.
- ✓ Ideal for microservices, replication, and data merging.
- ✓ UUIDs are non-sequential, improving security but affecting index locality.
- ✓ Store UUIDs using the `uuid` data type for best performance.

Best Practices

- Use UUID as PRIMARY KEY in distributed systems.
- Use `gen_random_uuid()` for better performance and simplicity.
- Avoid mixing UUID and BIGSERIAL in the same column for same purpose.
- Consider index impact due to randomness.
- Validate UUID format in application layer if needed.

Use Cases



Distributed
Systems



Data Merging
Across Systems



Microservices
Architecture



Data
Replication

Generating UUIDs

<code>gen_random_uuid()</code>	Generates a random UUID (requires <code>pgcrypto</code> extension)	<pre>SELECT gen_random_uuid(); 550e8400-e29b-41d4-a716-446655440000</pre>
<code>uuid_generate_v4()</code>	Generates a random (v4) UUID (requires <code>uuid-oss</code> extension)	<pre>SELECT uuid_generate_v4(); f81d4fae-7dec-11d0-a765-00a0c91e6bf6</pre>
<code>uuid_generate_v1mc()</code>	Generates time-based UUID (MAC address hidden)	<pre>SELECT uuid_generate_v1mc(); a4c13c10-8f2e-11ee-be56-0242ac120002</pre>

UUID Version Summary

Version	Name	Description	Use Case
v1	Time-based	Timestamp + MAC address	Legacy systems
v2	DCE Security	Time-based with POSIX UID	Rarely used
v3	Name-based (MD5)	Hash of name + namespace	Deterministic IDs
v4	Random	Random (122 bits of randomness)	Most common
v5	Name-based (SHA-1)	Hash of name + namespace	Deterministic IDs



Enable required extensions if needed: `CREATE EXTENSION IF NOT EXISTS "pgcrypto"; -- for gen_random_uuid()`

`CREATE EXTENSION IF NOT EXISTS "uuid-oss"; -- for uuid_generate_*`

JSON AND JSONB OVERVIEW

PostgreSQL can store and query semi-structured JSON directly in a relational table -- JSONB is almost always the right choice.

Type	Storage	When To Use
JSON	Stored as exact input text; re-parsed on every read.	Rare -- only when preserving exact formatting/whitespace matters.
JSONB	Stored in a decomposed binary form; indexable and queryable.	The default choice for any JSON column -- faster to query, supports indexing.

KEY TAKEAWAY JSONB stores JSON in an efficient, indexable binary form and supports operators like -> and @> for querying inside the document -- use it for genuinely flexible data, not as an escape hatch from schema design.

TRY IT YOURSELF — EXERCISE 1: CHOOSE COLUMN TYPES

DELIBERATELY

Reinforces: the PostgreSQL data type slides you just walked — numeric, character, boolean, date/time, UUID, and JSONB — applied to the CRM customer table.

STEPS (notes/lab37-type-choices.md)

Step	What To Do
1 — Column Table	For each customer column, name the PostgreSQL type you would use.
2 — Money Rule	Say which numeric type holds a monetary balance, and which one must never.
3 — Time Zones	Choose between timestamp and timestamptz, and give the reason.
4 — JSONB Limit	Say when a JSONB column is appropriate and when a real column is better.

Type choices

```
id uuid | text 'CUS-1001'  name text  email text  status text
balance numeric(12,2) -- never float/double for money
created_at timestamptz -- absolute instant, not a wall-clock guess
```

TRY IT NOW — Complete Exercise 1 before continuing: exercises/exercise-01-type-choices.md (workspace: examples/module-37-exercises).



PostgreSQL JSON and JSONB Overview

Store and query semi-structured data efficiently within PostgreSQL.
Choose between JSON (text) and JSONB (binary) based on your needs.

JSON vs JSONB Comparison

Feature	JSON (text)	JSONB (binary)
Storage Format	Stored as text	Stored in decomposed binary format
Parsing	Parsed on each read	Parsed on write, stored in binary
Performance	Slower for queries	Faster for queries and indexing
Indexing Support	Limited	Full GIN/GiST indexing support
Key Order	Preserves key order	Does not preserve key order
Duplicate Keys	Allowed	Last key wins (duplicates removed)
Whitespace	Preserved	Removed
Storage Size	Larger	Smaller (more efficient)
Use Case	Write-once, read-rarely, raw JSON	Read/Write worloarely, querying, indexing

Recommendation: Use JSONB for most use cases due to better performance, indexing, and flexibility. Use JSON only when you need to preserve exact text and formatting.

Example JSON Document

```
{
  "id": 101,
  "name": "Alice Johnson",
  "email": "alice.johnson@example.com",
  "active": true,
  "roles": ["admin", "user"],
  "address": {
    "street": "123 Main St",
    "city": "Dallas",
    "state": "TX",
    "zip": "75001"
  },
  "preferences": {
    "theme": "dark",
    "notifications": true
  },
  "createdAt": "2025-05-18T14:30:00Z"
}
```

Key Takeaways

- ✓ JSON stores data as plain text; JSONB stores in binary format.
- ✓ JSONB is faster for querying, indexing, and filtering.
- ✓ JSONB supports GIN indexes and advanced operators.
- ✓ JSON preserves formatting, order, and duplicate keys.
- ✓ Use JSONB for most production applications.
- ✓ Both support flexible, schema-less data modeling.
- ✓ Combine JSONB with relational data for hybrid models.

Common Use Cases



User Profiles
and Preferences



Product Catalogs
and Attributes



IoT and Sensor
Data



Audit Logs
and Events



Config and
Metadata

Creating Table with JSONB

```
CREATE TABLE employee (
  id          SERIAL PRIMARY KEY,
  name        TEXT NOT NULL,
  details     JSONB,
  created_at  TIMESTAMP DEFAULT now()
);
```

Use JSONB for indexing and efficient querying.

Querying JSONB Data (Examples)

Get value by key	<code>SELECT details->>'name' AS name FROM employee;</code>
Get nested value	<code>SELECT details->'address'->>'city' AS city FROM employee;</code>
Filter by key value	<code>SELECT * FROM employee WHERE details->>'active' = 'true';</code>
Check if key exists	<code>SELECT * FROM employee WHERE details ? 'roles';</code>
Containment	<code>SELECT * FROM employee WHERE details @> '{"roles": ["admin"]}';</code>
Array contains value	<code>SELECT * FROM employee WHERE details->'roles' ? 'admin';</code>

Indexing JSONB Data

```
-- Create GIN index for fast JSONB queries
CREATE INDEX idx_employee_details_gin .
ON employee USING GIN (details);
```

Useful operators supported by JSONB:

Operator	Description	?	Key exists
->	Get JSON object field	?1	Any key exists in array
->>	Get JSON object field as text	?&	All keys exist in array
#>	Get JSON value by path	@>	Contains
#>>	Get JSON value by path as text	<@	Is contained by

Best Practice: Use JSONB for performance, indexing, and querying. Index frequently queried keys. Validate JSON structure at the application layer if needed.

Since PostgreSQL 9.2 (JSON) and 9.4 (JSONB)

CREATING DATABASES

CREATE DATABASE provisions a new, fully isolated database inside a PostgreSQL cluster.

```
CREATE DATABASE crmdb
  WITH OWNER = crm_admin
      ENCODING = 'UTF8';

-- list databases in psql
\l
```

Run CREATE DATABASE while connected to any existing database (commonly postgres).

Always specify ENCODING = 'UTF8' explicitly -- don't rely on the cluster default.

DROP DATABASE permanently deletes it and every schema, table, and row inside -- no undo.

One database per application is the common default, e.g. crmdb for Northstar CRM.

KEY TAKEAWAY CREATE DATABASE provisions a new, isolated database; run it once per application, always with an explicit UTF8 encoding, and treat DROP DATABASE as irreversible.



PostgreSQL: Creating Databases

A database is a container for schemas, tables, and other database objects.
Use `CREATE DATABASE` to create a new database in PostgreSQL.

Overview

Databases in PostgreSQL are independent of each other.
Each database has its own:

- Schemas
- Tables and objects
- Users and permissions
- Configuration settings
- Storage (files on disk)

Syntax

```
CREATE DATABASE database_name
[ WITH
  OWNER = owner_name
  ENCODING = 'UTF8'
  LC_COLLATE = 'en_US.UTF-8'
  LC_CTYPE = 'en_US.UTF-8'
  TABLESPACE = tablespace_name
  CONNECTION LIMIT = -1
  IS_TEMPLATE = template_name ];
```

List Existing Databases

List all databases in the cluster:

```
\l
```

Or using SQL:

```
SELECT datname, owner, encoding, datcollate, datctype
FROM pg_database
ORDER BY datname;
```



Databases are stored in the data directory on disk as separate folders. Deleting a database removes all its objects permanently.

Examples

1 Create a simple database

```
CREATE DATABASE employee_db;
```

Creates a database with default settings.

2 Create a database with owner

```
CREATE DATABASE employee_db
WITH OWNER = app_user;
```

Sets the owner of the database.

3 Create a database with encoding and locale

```
CREATE DATABASE reporting_db
WITH ENCODING = 'UTF8'
LC_COLLATE = 'en_US.UTF-8'
LC_CTYPE = 'en_US.UTF-8';
```

Specifies encoding and locale for the database.

4 Create a database with tablespace and connection limit

```
CREATE DATABASE analytics_db
WITH TABLESPACE = fastspace
CONNECTION LIMIT = 50;
```

Stores data in the specified tablespace and limits concurrent connections.

5 Create a database from a template

```
CREATE DATABASE dev_db
WITH TEMPLATE template0
OWNER = dev_user;
```

Creates a database using an existing template.



Connect to the New Database

Connect to the newly created database:

```
\c employee_db
```

Verify the current database:

```
SELECT current_database();
```



Prerequisites

- You must have the `CREATEDB` privilege.
- You must connect to a **database** other than the one you are creating.
- Use a **superuser** account or a user with required permissions.



Best Practices

- ✓ Use meaningful and consistent database names.
- ✓ Assign a dedicated owner for each database.
- ✓ Choose the correct encoding and locale based on application needs.
- ✓ Use appropriate tablespaces for performance and storage management.
- ✓ Set connection limits to protect the system from too many connections.
- ✓ Regularly back up important databases.



Common Templates

Template	Description
template0	Empty template with no objects or settings.
template1	Default template with standard objects and settings.
Custom Template	A copy of an existing database for consistent setup.



Note: Once created, you can manage schemas, tables, users, permissions, and other objects within the database.



Be careful when dropping databases – there is no UNDO.

CREATING SCHEMAS

CREATE SCHEMA adds a named namespace inside the current database for grouping related tables.

```
CREATE SCHEMA crm AUTHORIZATION crm_app;  
  
-- qualify a table with its schema  
SELECT * FROM crm.customer;  
  
-- or set a default search path  
SET search_path TO crm, public;
```

Every database starts with a default schema named public.

search_path controls which schema an unqualified table name resolves to.

Schemas group tables logically -- by module, feature area, or tenant.

For Lab 37, the default public schema is enough -- extra schemas are optional polish.

KEY TAKEAWAY A schema is a namespace inside a database, resolved through search_path -- most applications, including Lab 37, can stay entirely in the default public schema.



Creating Schemas

Schemas help organize database objects into logical groups within a database. The **public** schema is created by default, but you can create custom schemas to support modular, secure, and maintainable database design.



Logical Organization

Group related tables, views, functions, and other objects under a common namespace.



Security & Access Control

Grant permissions at the schema level to control access for users and roles.



Modularity

Isolate applications or business domains in separate schemas within the same database.




Manageability

Simplifies maintenance, deployment, and collaboration across teams.



Key Takeaway: Schemas provide structure, security, and scalability within a database. Use schemas to organize objects, enforce access control, and support modular application design.

Database → Schemas → Objects



SQL Example: Create a Schema

```
1 -- Create a new schema
2 CREATE SCHEMA hr;
3
4 -- Create a schema with authorization
5 CREATE SCHEMA finance AUTHORIZATION app_user;
6
7 -- List all schemas in the database
8 SELECT schema_name FROM information_schema.schemata;
```



The **public** schema is available by default.

You can set your **search_path** to control which schemas are used by default.

CREATING TABLES

CREATE TABLE defines a table's columns, types, and constraints in one statement.

```
CREATE TABLE customer (  
  customer_id BIGINT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
  full_name   VARCHAR(150) NOT NULL,  
  email       VARCHAR(255) NOT NULL UNIQUE,  
  status      VARCHAR(20)  NOT NULL DEFAULT 'PROSPECT',  
  created_at  TIMESTAMPTZ  NOT NULL DEFAULT now()  
);
```

KEY TAKEAWAY CREATE TABLE declares every column's name, type, and constraints together -- the primary key, NOT NULL, UNIQUE, and DEFAULT clauses shown here are each covered in depth next.



Creating Tables

Tables store data in rows and columns. A well-designed table structure ensures data integrity, performance, and scalability. Define columns, data types, and constraints to enforce business rules at the database level.



Structured Storage

Organize data in tables with rows (records) and columns (fields).



Data Integrity

Use constraints like **PRIMARY KEY**, **FOREIGN KEY**, **NOT NULL**, **UNIQUE**, and **CHECK** to maintain accurate and consistent data.



Performance

Choose appropriate data types and indexing strategies to optimize query performance.



Scalability & Maintainability


A clear table design supports future growth, easier maintenance, and team collaboration.



Key Takeaway: A well-designed table structure is the foundation of a strong database. It ensures data integrity, enforces business rules, and supports high-performance applications.

Example Table: employees

	Column Name	Data Type	Constraints	Description
🔑	employee_id	SERIAL	PRIMARY KEY	Unique identifier for employee
👤	first_name	VARCHAR(50)	NOT NULL	Employee first name
👤	last_name	VARCHAR(50)	NOT NULL	Employee last name
✉️	email	VARCHAR(100)	UNIQUE, NOT NULL	Employee email address
📅	hire_date	DATE	NOT NULL	Date of joining
💰	salary	NUMERIC(12,2)	CHECK (salary > 0)	Monthly salary
👥	department_id	INTEGER	FOREIGN KEY	Reference to departments table



SQL Example: Create Table

```
1 -- Create employees table
2 CREATE TABLE employees (
3     employee_id SERIAL PRIMARY KEY,
4     first_name VARCHAR(50) NOT NULL,
5     last_name VARCHAR(50) NOT NULL,
6     email VARCHAR(100) NOT NULL UNIQUE,
7     hire_date DATE NOT NULL,
8     salary NUMERIC(12,2) CHECK (salary > 0),
9     department_id INTEGER REFERENCES departments(department_id)
10 );
11
```



Best Practices

- Use meaningful and consistent table and column names.
- Choose the right data types.
- Define primary keys for every table.
- Apply constraints to enforce business rules.
- Index foreign keys for better performance.

PRIMARY KEYS

A primary key is a column, or set of columns, that uniquely identifies every row in a table.



Must Be Unique

No two rows can ever share the same primary key value.



Cannot Be NULL

Every row must have a real value -- PostgreSQL enforces this automatically.



Only One Per Table

A table can declare exactly one primary key, single or composite.



anchors Foreign Keys

Other tables reference this table safely by pointing at its primary key.

KEY TAKEAWAY A primary key is always unique and never NULL, and every foreign key relationship this module covers depends on one existing to point at.



Primary Keys



A **Primary Key** is a column or a set of columns that **uniquely identifies** each row in a table.

Key Characteristics



Uniqueness

Each value in the primary key must be unique.



NOT NULL

Primary key columns cannot contain NULL values.



Entity Integrity

Ensures each row can be reliably identified and referenced.



Referencing

Referenced by foreign keys in other tables to establish relationships.

Example Table: employees

employee_id (PK)	first_name	last_name	email	hire_date	salary
1	John	Doe	john.doe@example.com	2024-01-15	65000.00
2	Jane	Smith	jane.smith@example.com	2024-02-10	72000.00
3	Michael	Brown	michael.brown@example.com	2024-03-05	68000.00
4	Emily	Johnson	emily.johnson@example.com	2024-04-20	70000.00



Unique identifier for each employee



Duplicate values are not allowed

Why Primary Keys Matter



Guarantees row uniqueness



Enables reliable relationships



Improves data consistency



Optimizes query performance

SQL Examples

1. Define Primary Key Inline

```
1 CREATE TABLE employees (  
2   employee_id SERIAL PRIMARY KEY,  
3   first_name VARCHAR(50) NOT NULL,  
4   last_name VARCHAR(50) NOT NULL,  
5   email VARCHAR(100) UNIQUE,  
6   hire_date DATE NOT NULL  
7 );
```

2. Define Primary Key as a Table Constraint

```
1 CREATE TABLE employees (  
2   employee_id INT,  
3   first_name VARCHAR(50) NOT NULL,  
4   last_name VARCHAR(50) NOT NULL,  
5   email VARCHAR(100) UNIQUE,  
6   hire_date DATE NOT NULL,  
7   PRIMARY KEY (employee_id)  
8 );
```

Best Practices

- ✓ Use surrogate keys (e.g., SERIAL, UUID) instead of natural keys.
- ✓ Keep primary keys simple and immutable.
- ✓ Use meaningful names like id, user_id, order_id.
- ✓ Avoid updating primary key values.



Key Takeaway: Primary keys uniquely identify each row, enforce data integrity, and serve as the foundation for relationships and efficient database design.

FOREIGN KEYS

A foreign key is a column in one table that must reference an existing primary key value in another table.

```
CREATE TABLE account (  
    account_id    BIGINT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
    customer_id   BIGINT NOT NULL,  
    balance       NUMERIC(19,2) NOT NULL DEFAULT 0,  
    CONSTRAINT fk_account_customer  
        FOREIGN KEY (customer_id) REFERENCES customer(customer_id)  
);
```

Every customer_id in account must already exist in customer.customer_id, or be NULL.
Inserting account with a non-existent customer_id is rejected -- no orphan rows.
A foreign key can repeat (many accounts per customer) -- uniqueness is not required.

KEY TAKEAWAY A foreign key enforces referential integrity automatically -- account.customer_id can only hold a value that genuinely exists in customer.customer_id, or NULL.



Foreign Keys



A **Foreign Key** is a column (or set of columns) in one table that refers to the **Primary Key** of another table.

Purpose

Foreign keys enforce **referential integrity** by ensuring that related data between tables remains consistent.

Key Benefits



Data Integrity

Prevents orphan records and maintains valid relationships.



Consistency

Ensures that related data exists before it is referenced.



Relationship Management


Supports one-to-many, many-to-many, and other relationships.



Reliable Operations

Simplifies joins and improves trust in data accuracy.

Example Relationship: Employees and Departments



SQL Example: Creating a Foreign Key

```
1 -- Parent table: departments
2 CREATE TABLE departments (
3     department_id SERIAL PRIMARY KEY,
4     department_name VARCHAR(100) NOT NULL,
5     location VARCHAR(100)
6 );
7
8 -- Child table: employees with foreign key
9 CREATE TABLE employees (
10    employee_id SERIAL PRIMARY KEY,
11    first_name VARCHAR(50) NOT NULL,
12    last_name VARCHAR(50) NOT NULL,
13    email VARCHAR(100) UNIQUE,
14    department_id INT NOT NULL,
15    CONSTRAINT fk_employees_department
16        FOREIGN KEY (department_id)
17        REFERENCES departments(department_id)
18        ON UPDATE CASCADE
19        ON DELETE RESTRICT
20 );
```

Common Actions

ON DELETE CASCADE
Automatically deletes child rows when parent is deleted.

ON DELETE RESTRICT
Prevents deletion of parent if child rows exist.

ON UPDATE CASCADE
Updates child rows when parent key is updated.

Sample Data

departments

department_id (PK)	department_name	location
1	Human Resources	New York
2	Engineering	Austin
3	Finance	Chicago

employees

employee_id (PK)	first_name	last_name	department_id (FK)
101	John	Doe	2
102	Jane	Smith	2
103	Emily	Johnson	1
104	Michael	Brown	3



If a department is deleted, employees referencing that department cannot remain (depending on the **ON DELETE** rule).



Best Practices

- ✓ Always define foreign keys to enforce relationships.
- ✓ Choose appropriate **ON DELETE** and **ON UPDATE** actions.
- ✓ Index foreign key columns for better performance.
- ✓ Use meaningful constraint names.
- ✓ Design relationships carefully before enforcing constraints.



Key Takeaway: Foreign keys connect tables and enforce referential integrity. They ensure that relationships between data remain valid, consistent, and reliable.

UNIQUE CONSTRAINTS

A UNIQUE constraint ensures every value in a column, or set of columns, is distinct across the whole table.

```
email VARCHAR(255) NOT NULL UNIQUE,  
CONSTRAINT uk_customer_email UNIQUE (email)
```

A table can have many UNIQUE constraints, but only one PRIMARY KEY.

A nullable UNIQUE column allows at most one NULL in most configurations.

A duplicate insert fails with: duplicate key value violates unique constraint.

Typical candidates: email, username, and any other real-world 'must be distinct' value.

KEY TAKEAWAY UNIQUE guarantees distinct values but, unlike a primary key, a table can have several UNIQUE constraints and the column may still allow NULL.



Unique Constraints



A **UNIQUE** constraint ensures that all values in a column (or set of columns) are distinct across all rows in a table.

Key Benefits



Prevents Duplicate Data

Ensures no two rows have the same value(s) in the constrained column(s).



Enforces Business Rules

Useful for columns like email, username, code, or any attribute that must be distinct.



Improves Data Quality

Helps maintain accuracy and reliability of database records.



Supports Performance

Unique constraints are automatically indexed by PostgreSQL for fast lookups.



Key Takeaway: Unique constraints ensure that data remains distinct and reliable. Use them to enforce business rules, prevent duplicates, and improve overall data integrity.

Example: Employees Table

Column	Data Type	Constraints	Description
employee_id	SERIAL	PRIMARY KEY	Unique identifier for employee
first_name	VARCHAR(50)	NOT NULL	Employee first name
last_name	VARCHAR(50)	NOT NULL	Employee last name
email	VARCHAR(100)	UNIQUE, NOT NULL	Employee email (must be unique)
phone_number	VARCHAR(20)	UNIQUE	Employee phone (must be unique)
hire_date	DATE	NOT NULL	Date of joining

SQL Examples

```
1 -- 1. Defining a UNIQUE constraint inline
2 CREATE TABLE employees (
3     employee_id SERIAL PRIMARY KEY,
4     first_name VARCHAR(50) NOT NULL,
5     last_name VARCHAR(50) NOT NULL,
6     email VARCHAR(100) NOT NULL UNIQUE,
7     phone_number VARCHAR(20) UNIQUE,
8     hire_date DATE NOT NULL
9 );
10 -- 2. Adding a UNIQUE constraint to an existing table
11 ALTER TABLE employees
12 ADD CONSTRAINT uq_employees_email UNIQUE (email);
```

Notes

- Multiple **UNIQUE** constraints can exist in a table.
- **UNIQUE** allows one **NULL** value (or multiple **NULLs** in PostgreSQL).
- For multiple columns, use a combined **UNIQUE** constraint.

Sample Data (emails must be unique)

employee_id	first_name	email
1	John	john.doe@example.com
2	Jane	jane.smith@example.com
3	Michael	michael.brown@example.com
4	Emily	emily.johnson@example.com



Violation Example

Inserting another row with email = jane.smith@example.com will fail.



Best Practices

- ✓ Use meaningful, business-specific column names.
- ✓ Apply **UNIQUE** constraints to natural keys (e.g., email, code).
- ✓ Prefer **UNIQUE** over application-level checks.
- ✓ Consider case-insensitive uniqueness (using **citext** extension if needed).
- ✓ Combine with **NOT NULL** unless **NULLs** are intentionally allowed.

NOT NULL CONSTRAINTS

A NOT NULL constraint ensures a column can never hold NULL -- it makes a value mandatory.

Column	Constraint	Why
full_name	NOT NULL	A customer record without a name is not usable.
email	NOT NULL	Required for every notification and login flow.
middle_name	nullable (no NOT NULL)	Genuinely optional -- not every customer has one.
closed_at	nullable (no NOT NULL)	Only set once an account is actually closed.

KEY TAKEAWAY Default to NOT NULL for every column a business rule truly requires, and leave a column nullable only when the absence of a value is itself a meaningful, expected state.



NOT NULL Constraints



A **NOT NULL** constraint prevents a column from storing **NULL** values. It ensures that every row must have a value in the constrained column(s).

Key Benefits



Ensures Mandatory Data

Guarantees that important columns always contain a value.



Improves Data Integrity

Reduces incomplete or missing information in the database.



Supports Business Rules

Enforces rules that certain fields are required for valid records.



Better Application Reliability

Applications can rely on the presence of data without additional **NULL** checks.

Example: Employees Table

Column Name	Data Type	Constraints	Description
employee_id	SERIAL	PRIMARY KEY	Unique identifier
first_name	VARCHAR(50)	NOT NULL	Employee first name (required)
last_name	VARCHAR(50)	NOT NULL	Employee last name (required)
email	VARCHAR(100)	UNIQUE, NOT NULL	Employee email (required & unique)
phone_number	VARCHAR(20)	NULL	Employee phone (optional)
hire_date	DATE	NOT NULL	Date of joining (required)
salary	NUMERIC(12,2)	NOT NULL	Salary (required)

Sample Data (NOT NULL enforced)

employee_id	first_name	last_name	email
1	John	Doe	john.doe@example.com
2	Jane	Smith	jane.smith@example.com
3	Michael	Brown	michael.brown@example.com



What Happens if NOT NULL is Violated?

Inserting a row with **NULL** in a **NOT NULL** column will result in an error.

ERROR: null value in column "email" of relation "employees" violates not-null constraint

SQL Examples

1. Defining NOT NULL in Table Creation

```
1 CREATE TABLE employees (  
2   employee_id SERIAL PRIMARY KEY,  
3   first_name  VARCHAR(50) NOT NULL,  
4   last_name   VARCHAR(50) NOT NULL,  
5   email       VARCHAR(100) NOT NULL,  
6   hire_date   DATE NOT NULL,  
7   salary      NUMERIC(12,2) NOT NULL,  
8   phone_number VARCHAR(20)  
9 );
```

2. Adding NOT NULL to Existing Column

```
1 ALTER TABLE employees  
2 ALTER COLUMN hire_date SET NOT NULL;
```



You cannot add **NOT NULL** if any existing row contains **NULL** in that column. Update or remove **NULLs** before applying.



Best Practices

- ✓ Apply **NOT NULL** to all required business-critical fields.
- ✓ Use **NOT NULL** along with appropriate data types, constraints, and defaults.
- ✓ Avoid allowing **NULLs** unless the data is truly optional.
- ✓ Validate data at the application level to prevent constraint violations.
- ✓ Review and clean existing data before adding **NOT NULL** constraints.



Key Takeaway: **NOT NULL** constraints ensure that columns always contain meaningful data. They help maintain data quality, enforce business rules, and make applications more reliable.

CHECK CONSTRAINTS

A CHECK constraint restricts the values a column can hold by enforcing a boolean condition.

```
CREATE TABLE customer (  
    customer_id BIGINT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
    status      VARCHAR(20) NOT NULL DEFAULT 'PROSPECT',  
    CONSTRAINT ck_customer_status  
        CHECK (status IN ('PROSPECT', 'ACTIVE', 'SUSPENDED', 'CLOSED'))  
);
```

KEY TAKEAWAY A CHECK constraint enforces a business rule directly in the database -- an insert or update that violates it is rejected before the row ever reaches storage.



CHECK Constraints

A **CHECK** constraint ensures that all values in a column (or set of columns) satisfy a specific condition or rule.

Key Benefits

- Enforces Business Rules**
Validates data based on application logic (e.g., salary must be greater than 0).
- Improves Data Integrity**
Prevents invalid or out-of-range data from being stored.
- Adds Flexibility**
Allows complex validations beyond data types, NOT NULL, or UNIQUE constraints.
- Catches Errors Early**
Ensures only valid data enters the database, reducing downstream issues.

Key Takeaway: CHECK constraints enforce domain-specific business rules and ensure that only valid, meaningful data is stored in your PostgreSQL database.

Example: Employees Table with CHECK Constraints

Column Name	Data Type	Constraints	Description	Example Rule
employee_id	SERIAL	PRIMARY KEY	Unique identifier	–
first_name	VARCHAR(50)	NOT NULL	Employee first name	–
last_name	VARCHAR(50)	NOT NULL	Employee last name	–
email	VARCHAR(100)	UNIQUE, NOT NULL	Employee email	–
salary	NUMERIC(12,2)	NOT NULL, CHECK (salary > 0)	Salary must be greater than 0	salary > 0
hire_date	DATE	NOT NULL, CHECK (hire_date <= CURRENT_DATE)	Hire date cannot be in the future	hire_date <= CURRENT_DATE
department_id	INT	NOT NULL, CHECK (department_id > 0)	Department id must be a positive number	department_id > 0
status	VARCHAR(20)	NOT NULL, CHECK (status IN ('ACTIVE','INACTIVE'))	Status must be one of the allowed values	status IN ('ACTIVE','INACTIVE')

SQL Examples

1. Creating Table with CHECK Constraints

```
1 CREATE TABLE employees (  
2   employee_id SERIAL PRIMARY KEY,  
3   first_name  VARCHAR(50) NOT NULL,  
4   last_name   VARCHAR(50) NOT NULL,  
5   email       VARCHAR(100) UNIQUE NOT NULL,  
6   salary      NUMERIC(12,2) NOT NULL  
7   CHECK (salary > 0),  
8   hire_date   DATE NOT NULL  
9   CHECK (hire_date <= CURRENT_DATE),  
10  department_id INT NOT NULL  
11  CHECK (department_id > 0),  
12  status      VARCHAR(20) NOT NULL  
13  CHECK (status IN ('ACTIVE','INACTIVE'))  
14 );
```

2. Adding a CHECK Constraint to Existing Table

```
1 ALTER TABLE employees  
2 ADD CONSTRAINT chk_salary_positive  
  CHECK (salary > 0);
```

3. Dropping a CHECK Constraint

```
1 ALTER TABLE employees  
2 DROP CONSTRAINT chk_salary_positive;
```

CHECK constraints are enforced per row and cannot reference other rows. They help keep data accurate and consistent based on defined business rules.

Sample Data

employee_id	salary	hire_date	status
1	75000.00	2023-05-10	ACTIVE
2	62000.00	2022-11-01	ACTIVE
3	58000.00	2024-03-20	INACTIVE

Constraint Violation Examples

- salary = -5000 → Violates: salary > 0
- hire_date = '2026-01-01' → Violates: hire_date <= CURRENT_DATE
- status = 'PENDING' → Violates: status IN ('ACTIVE','INACTIVE')
- department_id = 0 → Violates: department_id > 0

Best Practices

- ✓ Keep CHECK conditions simple and clear.
- ✓ Use meaningful names for constraints.
- ✓ Validate rules at the database level to protect data integrity.
- ✓ Test constraints with valid and invalid data.
- ✓ Avoid heavy computations or subqueries in CHECK conditions.

DEFAULT VALUES

A DEFAULT clause supplies a column's value automatically whenever an INSERT does not specify one.

Column	Default	Behavior
status	DEFAULT 'PROSPECT'	New customers start as PROSPECT unless stated otherwise.
created_at	DEFAULT now()	Stamped with the current instant automatically.
balance	DEFAULT 0	New accounts start at zero, never left NULL.
customer_id	GENERATED ALWAYS AS IDENTITY	A generated value, conceptually similar to a default.

KEY TAKEAWAY A DEFAULT is applied only when a column is omitted from an INSERT entirely -- it never overrides a value the statement actually supplies, including an explicit NULL.

TRY IT YOURSELF — EXERCISE 2: PLAN KEYS AND CONSTRAINTS

Reinforces: creating tables, primary and foreign keys, unique, not null, check constraints, and defaults as you just covered them.

STEPS (notes/lab37-constraints.md)

Step	What To Do
1 — Key Choice	Name the primary key for customer and account, and say why.
2 — Constraint List	Give one unique, one not-null, and one check constraint for the CRM.
3 — Foreign Key	Write the account-to-customer foreign key and its delete behaviour.
4 — Why in the Database	Say why these rules belong in the schema and not only in Java.

CRM constraints

```
customer(id pk)  account(id pk, customer_id references customer(id))
unique(email)   not null(name, status)
check (status in ('ACTIVE','PROSPECT'))  -- the database is the last line
```

TRY IT NOW — Complete Exercise 2 before continuing: exercises/exercise-02-constraint-plan.md (workspace: examples/module-37-exercises).



DEFAULT Values



A **DEFAULT** value is used when no value is provided for a column during INSERT. The database automatically assigns the default value defined for that column.

Key Benefits



Ensures Consistency
Applies a standard value when none is provided, keeping data consistent.



Reduces Errors
Prevents NULL or missing values for important columns.



Improves Productivity
Simplifies INSERT operations by reducing the need to specify values for every column.



Supports Business Rules
Helps enforce default business logic (e.g., status = 'ACTIVE').



Key Takeaway: DEFAULT values help ensure data completeness, maintain consistency, and simplify data entry by automatically providing values when none are supplied.

Example: Employees Table with DEFAULT Values

Column Name	Data Type	Constraints / Default	Description	Example Default
employee_id	SERIAL	PRIMARY KEY	Unique identifier	-
first_name	VARCHAR(50)	NOT NULL	Employee first name	-
last_name	VARCHAR(50)	NOT NULL	Employee last name	-
email	VARCHAR(100)	UNIQUE, NOT NULL	Employee email	-
status	VARCHAR(20)	NOT NULL DEFAULT 'ACTIVE'	Employment status	'ACTIVE'
hire_date	DATE	NOT NULL DEFAULT CURRENT_DATE	Date of joining	Current date
created_at	TIMESTAMP	NOT NULL DEFAULT CURRENT_TIMESTAMP	Record creation time	Current timestamp
salary	NUMERIC(12,2)	NOT NULL DEFAULT 0	Employee salary	0

SQL Examples

1. Creating Table with DEFAULT Values

```
1 CREATE TABLE employees (  
2   employee_id SERIAL PRIMARY KEY,  
3   first_name VARCHAR(50) NOT NULL,  
4   last_name VARCHAR(50) NOT NULL,  
5   email VARCHAR(100) UNIQUE NOT NULL,  
6   status VARCHAR(20) NOT NULL DEFAULT 'ACTIVE',  
7   hire_date DATE NOT NULL DEFAULT CURRENT_DATE,  
8   created_at TIMESTAMP NOT NULL  
9   DEFAULT CURRENT_TIMESTAMP,  
10  salary NUMERIC(12,2) NOT NULL DEFAULT 0  
11 );  
12
```

2. Altering Table to Add a DEFAULT

```
1 ALTER TABLE employees  
2 ALTER COLUMN status SET DEFAULT 'ACTIVE';  
  
4 ALTER TABLE employees  
5 ALTER COLUMN salary SET DEFAULT 0;
```



DEFAULT values are applied only when a value is not provided in the INSERT statement.

Sample Inserts and Results

1. Insert without specifying default columns

```
INSERT INTO employees (first_name, last_name, email)  
VALUES ('John', 'Doe', 'john.doe@example.com');
```



status	hire_date	created_at	salary
'ACTIVE'	2025-05-11	2025-05-11 10:30:45	0.00

2. Insert with custom values (overrides defaults)

```
INSERT INTO employees (first_name, last_name, email,  
status, salary)  
VALUES ('Jane', 'Smith', 'jane.smith@example.com',  
'INACTIVE', 75000.00);
```



status	hire_date	created_at	salary
'INACTIVE'	2025-05-11	2025-05-11 10:30:46	75000.00



Best Practices

- ✓ Use DEFAULT for values that are common or expected.
- ✓ Use meaningful defaults (e.g., CURRENT_DATE, 0, 'ACTIVE').
- ✓ Avoid hiding data issues behind defaults.
- ✓ Review defaults carefully before applying to production.
- ✓ Document default values for clarity.

TABLE RELATIONSHIPS

A relationship defines how two or more tables are associated -- the core of relational schema design.

Cardinality	Notation	Enforced By
One-to-One	1 : 1	Foreign key with a UNIQUE constraint on the child.
One-to-Many	1 : M	Foreign key on the 'many' side referencing the 'one' side.
Many-to-Many	M : M	A junction table with foreign keys to both sides.

KEY TAKEAWAY Every relationship in a relational schema reduces to one of three cardinalities, and each is implemented with foreign keys -- the next slide walks a worked example of each.


Table Relationships

Relationships define how data in one table connects to data in another.



1 One-to-One

A record in Table A relates to one and only one record in Table B.




`users.id = user_profiles.user_id`

- Used when additional details are stored separately.
- Example: users and user_profiles
- Implement with UNIQUE foreign key in one of the tables.

2 One-to-Many

A record in Table A relates to many records in Table B.




`employees.department_id → departments.id`

- Most common relationship.
- Example: departments and employees
- Foreign key in the "many" table references the "one" table.

3 Many-to-One

Many records in Table A relate to one record in Table B.




`orders.customer_id → customers.id`

- A variation of one-to-many.
- Example: orders and customers
- Many orders belong to one customer.

4 Many-to-Many

Many records in Table A relate to many records in Table B.



- Requires a join table.
- Example: students and courses
- Join table contains foreign keys to both tables.

Key Points

Relationships enforce data integrity.

Use foreign keys to maintain valid references.

Proper relationships improve query performance.

Choose the right relationship to match business needs.

Well-designed relationships support scalability.



KEY TAKEAWAY: Table relationships model real-world connections, enforce data integrity, and enable powerful queries.

ONE-TO-ONE, ONE-TO-MANY, MANY-TO-MANY

Three worked examples from the Northstar CRM schema -- one for each cardinality.

1

One-to-One

customer <-> customer_profile. A UNIQUE foreign key on customer_profile.customer_id guarantees at most one profile per customer.



One-to-Many

customer -> account. A plain foreign key, account.customer_id, lets one customer own many accounts.




Many-to-Many

account <-> tag, resolved by a junction table account_tag(account_id, tag_id) with a foreign key to each side.

KEY TAKEAWAY One-to-one adds UNIQUE to the foreign key, one-to-many uses a plain foreign key on the many side, and many-to-many always needs a junction table -- there is no fourth pattern.


One-to-One

A record in Table A relates to one and only one record in Table B, and vice versa.



EXAMPLE

Each user has exactly one user profile, and each user profile belongs to exactly one user.



users.id = user_profiles.user_id

The foreign key in user_profiles is also the primary key, ensuring one-to-one relationship.

KEY CHARACTERISTICS

- 1:1** One record in Table A relates to one record in Table B.
- Bidirectional** The relationship is mutual.
- Enforces Uniqueness** The foreign key must be UNIQUE (or the primary key).
- Used for Separate Details** Useful when additional details are stored in a separate table.
- Integrity Guaranteed** Deleting a record in one table affects the related record in the other.

SQL EXAMPLE

Create users table

```
CREATE TABLE users (
  id SERIAL PRIMARY KEY,
  username VARCHAR(50) NOT NULL,
  email VARCHAR(100) UNIQUE NOT NULL,
  created_at TIMESTAMP DEFAULT NOW()
);
```

Create user_profiles table

```
CREATE TABLE user_profiles (
  user_id INTEGER PRIMARY KEY
    REFERENCES users(id)
    ON DELETE CASCADE,
  full_name VARCHAR(100) NOT NULL,
  phone VARCHAR(20),
  address VARCHAR(255),
  date_of_birth DATE
);
```




KEY TAKEAWAY: One-to-one relationships are ideal when data is logically separated but tightly linked, ensuring each record in one table has exactly one matching record in the other.




One-to-Many

A single record in Table A relates to many records in Table B, but each record in Table B relates to only one record in Table A.



EXAMPLE

One department can have many employees.
Each employee belongs to one department.



employees.department_id is a foreign key that references **departments.id**.
It ensures every employee belongs to an existing department.

KEY CHARACTERISTICS

1:N

One-to-many cardinality

One record in Table A → many records in Table B.



Foreign Key on Many Side

The “many” table contains a foreign key referencing the “one” table.



Referential Integrity

A department must exist before an employee can reference it.



Supports Business Hierarchies

Commonly used for master-detail relationships.



Cascade Options

ON DELETE CASCADE can be used to remove employees with a department.

SQL EXAMPLE

Create departments table

```
CREATE TABLE departments (
  id SERIAL PRIMARY KEY,
  name VARCHAR(100) NOT NULL,
  location VARCHAR(100)
);
```

Create employees table

```
CREATE TABLE employees (
  id SERIAL PRIMARY KEY,
  name VARCHAR(100) NOT NULL,
  email VARCHAR(100) NOT NULL,
  department_id INTEGER NOT NULL,
  CONSTRAINT fk_employees_department
  FOREIGN KEY (department_id)
  REFERENCES departments(id)
  ON DELETE RESTRICT
);
```



KEY TAKEAWAY: Use one-to-many relationships to model hierarchical or master-detail data.
They are the most common relationship type in relational databases.



One
Department




Many
Employees

Many-to-Many

Many records in Table A relate to many records in Table B, and many records in Table B relate to many records in Table A.

EXAMPLE

A student can enroll in many courses.
A course can have many students.



student_courses is a junction (bridge) table.

It contains foreign keys from both tables.

The combination of **student_id** and **course_id** is the primary key.

KEY CHARACTERISTICS

- M:N** **Many-to-many cardinality**
Many records in Table A relate to many records in Table B.
- Requires a Junction Table**
A third table stores the relationships using foreign keys from both tables.
- Composite Primary Key**
The junction table uses a composite primary key (student_id, course_id) to prevent duplicate relationships.
- Additional Attributes**
The junction table can store extra information about the relationship, such as enrollment date or grade.
- Supports Scalability**
Efficient and flexible for real-world relationships.

SQL EXAMPLE

Create students table

```
CREATE TABLE students(
  id SERIAL PRIMARY KEY,
  first_name VARCHAR(50) NOT NULL,
  last_name VARCHAR(50) NOT NULL,
  email VARCHAR(100) UNIQUE
);
```

Create courses table


```
CREATE TABLE courses (
  id SERIAL PRIMARY KEY,
  title VARCHAR(100) NOT NULL,
  credits INTEGER NOT NULL,
  description TEXT
);
```

Create junction table

```
CREATE TABLE student_courses (
  student_id INTEGER NOT NULL REFERENCES students(id),
  course_id INTEGER NOT NULL REFERENCES courses(id),
  enrolled_at TIMESTAMP DEFAULT NOW(),
  grade VARCHAR(10),
  PRIMARY KEY (student_id, course_id)
);
```



KEY TAKEAWAY: Use many-to-many relationships when both sides can have multiple matches. Model them using a junction table with foreign keys to both tables.



NORMALIZATION OVERVIEW

Normalization organizes data to reduce redundancy and eliminate update, insert, and delete anomalies.



Update Anomaly

The same fact repeated in many rows -- update one copy and the others silently go stale.



Insert Anomaly

Cannot record one fact without also inventing unrelated, not-yet-known data.



Delete Anomaly

Deleting one row accidentally destroys other, still-needed information.



The Fix

Decompose overloaded tables into smaller, well-structured tables connected by foreign keys.

KEY TAKEAWAY Normalization eliminates redundancy by giving every fact exactly one home -- most business schemas, including Lab 37's, target Third Normal Form (3NF).

TRY IT YOURSELF — EXERCISE 3: MODEL THE CRM RELATIONSHIPS

Reinforces: table relationships, the one-to-one, one-to-many, and many-to-many slides, and normalization as you just covered them.

STEPS (notes/lab37-er-notes.md)

Step	What To Do
1 — Cardinalities	State the cardinality between customer and account, in words.
2 — Many-to-Many	Name a CRM case needing a join table, and the columns that table holds.
3 — Normalize	Give one repeated value you would move into its own table, and why.
4 — When Not To	Name one case where duplicating a value is the right call.

CRM cardinalities

```
customer 1 --- * account      (one customer, many accounts)
customer * --- * tag via customer_tag(customer_id, tag_id)
denormalize deliberately: store the price as charged, not as looked up
```

TRY IT NOW — Complete Exercise 3 before continuing: `exercises/exercise-03-relationship-model.md` (workspace: `examples/module-37-exercises`).

Normalization Overview

Normalization organizes data to reduce redundancy and improve data integrity.

Denormalized				Normalized				
id	student	course	instructor	student_id	student	course_id	course	instructor
1	Alice	Java	Smith	1	Alice	101	Java	Smith
2	Bob	Angular	Jones	2	Bob	102	Angular	Jones
3	Alice	Spring	Smith	3	Alice	103	Spring	Smith

What is Normalization?



Normalization is the process of structuring a relational database to reduce data redundancy and improve data integrity by organizing data into multiple related tables.

Key Goals

- Eliminate duplicate data
- Maintain data consistency
- Simplify updates and maintenance
- Improve query performance

The Normal Forms

1NF

First Normal Form

Eliminate repeating groups and ensure atomic values in each column.

2NF

Second Normal Form

Remove partial dependencies (all non-key attributes depend on the whole key).

3NF

Third Normal Form

Remove transitive dependencies (non-key attributes depend only on the key).

BCNF

Boyce-Codd Normal Form

Stronger version of 3NF (all determinants must be candidate keys).

4NF

Fourth Normal Form

Remove multi-valued dependencies.

5NF

Fifth Normal Form

Remove join dependencies.

Example: From 1 Table to Normalized Tables

Before Normalization (1NF)

student_id	student_name	course_id	course_name	instructor
1	Alice	101	Java	Smith
1	Alice	102	Angular	Jones
2	Bob	103	Spring	Smith



After Normalization (3NF)

Students		Courses			Enrollments	
student_id	student_name	course_id	course_name	instructor	student_id	course_id
1	Alice	101	Java	Smith	1	101
2	Bob	102	Angular	Jones	1	102
		103	Spring	Smith	2	103



Key Takeaways

- Normalization reduces redundancy and maintains data integrity.
- Each normal form removes specific types of dependencies.
- Proper normalization leads to cleaner design and better performance.



Next Step

Design normalized tables for your domain model, starting with 1NF and progressing to 3NF.

INDEX FUNDAMENTALS

An index is a separate data structure that lets PostgreSQL find rows without scanning the whole table.



Book-Index Analogy

Like a book's index, it points straight to the right page instead of reading every page.



Speeds Up Reads

Turns a full table scan into a fast lookup for WHERE, JOIN, and ORDER BY.



Costs on Writes

Every INSERT, UPDATE, or DELETE must also maintain every index on that table.



Index With Intent

Add an index because a real query pattern needs it -- not on every column by default.

KEY TAKEAWAY An index trades write cost and storage for dramatically faster reads -- Module 38 goes deep on this trade-off, but the core idea starts here.

Index Fundamentals

Indexes improve the speed of data retrieval operations on a database table.

They work like the index in a book—helping the database find rows quickly without scanning the entire table.



Goal of Indexes

Reduce disk I/O, speed up searches and joins, and improve overall query performance.


How an Index Works (Example)

Table: **employees**

id	name	department	salary
1	Alice	Finance	70000
2	Bob	HR	60000
3	Carol	IT	80000
4	David	Finance	65000
5	Eve	IT	90000
6	Frank	HR	55000

Create Index
on department

Index on department



Query:

WHERE department = 'IT'



Database uses index
to quickly find rows

Row IDs: 3, 5

Key Points



Speeds Up Searches

Allows the database to locate rows without scanning the entire table.



Extra Storage

Indexes consume additional disk space and memory.



Slower Writes

INSERT, UPDATE, DELETE operations are slower because indexes must be maintained.



Balance is Key

Create indexes on columns used in WHERE, JOIN, ORDER BY, GROUP BY, but avoid over-indexing.

Common Index Types in PostgreSQL

B-Tree Index




Default index type.
Best for equality and range queries (<, <=, >, >=, BETWEEN).

Unique Index



Ensures all values in the indexed column(s) are unique.

Composite Index



Index on multiple columns.
Column order matters for query optimization.

Partial Index



Indexes only a subset of rows that meet a condition.
Smaller and more efficient.

Expression Index



Indexes the result of an expression or function.
Useful for computed values.

When to Create an Index

- ✓ Columns used in WHERE clauses
- ✓ Columns used in JOIN conditions
- ✓ Columns used in ORDER BY / GROUP BY
- ✓ Columns used in DISTINCT
- ✓ Columns with high selectivity (few unique values)
- ✗ Avoid columns with very low selectivity (e.g., Boolean, status flags)



Remember: Indexes are powerful but come with a cost.
Create the right indexes, not too many, and review them regularly.



PostgreSQL Tip

- Use EXPLAIN (ANALYZE, BUFFERS) to verify index usage.
- Drop unused indexes to save space and improve write performance.

B-TREE INDEXES

B-Tree is PostgreSQL's default index type -- efficient for equality, range, and sorted queries.

```
CREATE INDEX idx_customer_status ON customer (status);  
-- USING btree is the default and can be omitted
```

Works for =, <, <=, >, >=, BETWEEN, and IN comparisons.
Supports ORDER BY directly -- the index is already sorted.
The right default for most foreign keys and frequently filtered columns.

KEY TAKEAWAY CREATE INDEX defaults to a B-Tree, PostgreSQL's general-purpose index -- it is almost always the right first choice unless a query pattern calls for something more specialized.

B-Tree Indexes

B-Tree is the default index type in PostgreSQL. It is a balanced tree data structure that keeps index entries sorted, allowing efficient search, insert, update, and delete.



What is a B-Tree Index?


A B-Tree (Balanced Tree) index stores data in a sorted tree structure. Each node can have multiple keys and multiple children. All leaf nodes are at the same level, keeping the tree balanced for consistent performance.

How a B-Tree Index Works (Example)

Example Table: employees

id (PK)	name	department	salary
101	Alice	Finance	70000
102	Bob	HR	60000
103	Carol	IT	80000
104	David	Finance	65000
105	Eve	IT	90000
106	Frank	HR	55000
107	Grace	Finance	75000
108	Hannah	IT	68000
109	Ivy	HR	58000

B-Tree Index on id



Search Example:

Find id = 105

- 1 Start at root node [104]
105 > 104, go right
- 2 At node [108]
105 < 108, go left
- 3 At node [102, 106]
105 is between 102 and 106,
go to middle child
- 4 Reached leaf [105, 106]
105 found!

Key Characteristics



Balanced Tree Structure

The tree remains balanced after inserts and deletes, ensuring consistent performance.



Ordered Data

Index entries are stored in sorted order, enabling fast searches and range queries.



Efficient Operations

Optimized for equality searches and range-based queries.



Default Index Type

B-Tree is the default index type for most column data types in PostgreSQL.

Good For

- ✓ Equality comparisons (=)
- ✓ Range queries (>, >=, <, <=, BETWEEN)
- ✓ ORDER BY and GROUP BY operations
- ✓ JOIN operations

Not Ideal For

- ✗ Full-text search
- ✗ Leading wildcard searches (LIKE '%abc')
- ✗ Very low-selectivity columns (e.g., gender, status flags)

Advantages

- ✓ Fast search performance
- ✓ Handles large datasets efficiently
- ✓ Maintains performance as data grows
- ✓ Well-suited for most transactional workloads

Example in PostgreSQL

```
-- Create a B-Tree index on salary column
CREATE INDEX idx_employees_salary
ON employees (salary);
```



Key Takeaway: B-Tree indexes provide a balanced, sorted structure that accelerates data retrieval. They are the most commonly used indexes in relational databases and a foundation for query performance.



Tip: Monitor index usage with EXPLAIN (ANALYZE, BUFFERS) and remove unused indexes to save storage and improve write performance.

UNIQUE INDEXES

A UNIQUE index enforces distinct values while also speeding up lookups on that column.

```
CREATE UNIQUE INDEX idx_customer_email ON customer (email);
```

Behaves exactly like a UNIQUE constraint -- in fact, UNIQUE constraints are implemented as unique indexes internally.

A duplicate insert is rejected with the same 'violates unique constraint' error.

PRIMARY KEY and UNIQUE columns already get a unique index automatically -- no separate CREATE UNIQUE INDEX is needed for those.

KEY TAKEAWAY PostgreSQL implements every UNIQUE constraint, and every PRIMARY KEY, as a unique index under the hood -- CREATE UNIQUE INDEX is the same mechanism used explicitly.

Unique Indexes

A **Unique Index** enforces uniqueness of values in one or more columns.

It ensures that no two rows in a table can have the same value(s) in the indexed column(s).



Why Use Unique Indexes?

- ✓ Enforce business rules and data integrity
- ✓ Prevent duplicate data
- ✓ Improve query performance for equality searches
- ✓ Often used to support PRIMARY KEY and UNIQUE constraints

Example


Table: **users**

id (PK)	username	email	role
1	alice	alice@example.com	ADMIN
2	bob	bob@example.com	USER
3	carol	carol@example.com	USER

Unique Index on email

```
CREATE UNIQUE INDEX ux_users_email  
ON users (email);
```

How It Works



The index stores **unique** email values in sorted order.

Each value points to the corresponding row in the table.

Key Characteristics



Enforces Uniqueness

Prevents duplicate values in the indexed column(s).



Allows NULL Values

PostgreSQL allows multiple NULLs in a unique index (NULLs are not considered equal).



Improves Performance

Speeds up equality lookups and joins on the indexed column(s).



Can Be Composite

Unique indexes can be created on multiple columns to enforce combined uniqueness.

Enforcement Example

Valid Insert

```
INSERT INTO users (username, email, role)  
VALUES ('dave', 'dave@example.com', 'USER');
```



Invalid Insert (Duplicate Email)

```
INSERT INTO users (username, email, role)  
VALUES ('david', 'alice@example.com', 'USER');
```



Error Returned

ERROR: duplicate key value violates unique constraint "ux_users_email"

DETAIL: Key (email)=(alice@example.com) already exists.

Unique Index vs Other Indexes

Index Type	Allows Duplicates?	Enforces Uniqueness?	Typical Use Case
B-Tree Index	Yes	No	General search and performance
Unique Index	No	Yes	Enforce uniqueness, integrity
Primary Key	No	Yes	Entity identity (not null + unique)

Best Practices

- ✓ Create unique indexes to enforce business rules.
- ✓ Use meaningful column(s) with high selectivity.
- ✓ Avoid unnecessary unique indexes (write overhead).
- ✓ Use composite unique indexes for multi-column rules.
- ✓ Monitor index usage with EXPLAIN and pg_stat_user_indexes.



Key Takeaway: Unique indexes are essential for maintaining data integrity by ensuring that values in one or more columns remain unique. They also improve query performance for equality lookups while enforcing important business rules.



Tip: Use UNIQUE constraints or unique indexes based on whether the rule is part of the schema definition or created later for performance.

COMPOSITE INDEXES

A composite index spans two or more columns, sorted together in the order they're declared.

```
CREATE INDEX idx_account_customer_status ON account (customer_id, status);
```

Speeds up queries that filter on the leading column, or the leading column plus more.
Column order matters -- (customer_id, status) is not the same index as (status, customer_id).
Module 38 covers choosing composite column order in depth, using real query patterns.

KEY TAKEAWAY A composite index sorts by its first column, then its second, and so on -- get the column order right and one index can serve several related queries.

TRY IT YOURSELF — EXERCISE 4: PLAN INDEXES FOR REAL QUERIES

Reinforces: index fundamentals, B-tree, unique, and composite indexes as you just covered them, tied to real CRM queries.

STEPS (notes/lab37-indexes.md)

Step	What To Do
1 — Query First	Write the two queries the CRM runs most, before naming any index.
2 — Index Per Query	Give the index that serves each, including column order for composites.
3 — Column Order	Say why the composite's column order matters.
4 — Cost	Say what every extra index costs on write.

Index plan

```
q1 find by email      -> unique index on customer(email)
q2 accounts for a customer, newest first
  -> index on account(customer_id, opened_at desc)
```

TRY IT NOW — Complete Exercise 4 before continuing: exercises/exercise-04-index-plan.md (workspace: examples/module-37-exercises).



Composite Indexes

A composite index is an index on **two or more columns**. It improves query performance when filtering or sorting by multiple columns.

Example Table: **orders**

order_id (PK)	customer_id	order_date	status	total_amount
1001	501	2025-05-01	NEW	125.00
1002	501	2025-05-02	SHIPPED	220.00
1003	502	2025-05-01	NEW	75.00
1004	503	2025-05-03	DELIVERED	320.00
1005	501	2025-05-03	CANCELLED	60.00


```
CREATE INDEX idx_orders_customer_date  
ON orders (customer_id, order_date);
```

This composite index is ordered as:

customer_id → **order_date**

The column order matters.
PostgreSQL can use the index from **left to right**.

How a Composite Index Works



Efficient for queries filtering by **customer_id** alone, or by **customer_id** and **order_date** together.



Best Practices

- Place the **most selective** column first.
- Match the **index column order** with your most common query patterns.
- Avoid unnecessary composite indexes; they increase **write cost**.
- Check index usage with **EXPLAIN ANALYZE**.



Rule of Thumb: Use composite indexes when queries filter by multiple columns together.



PostgreSQL

POSTGRESQL SEQUENCES

A sequence is a standalone database object that generates a strictly increasing series of numbers.

```
CREATE SEQUENCE account_number_seq START 100000 INCREMENT 1;  
  
SELECT nextval('account_number_seq'); -- 100000  
SELECT nextval('account_number_seq'); -- 100001
```

Sequences are safe under concurrency -- two simultaneous callers never receive the same value.
A gap can appear if a transaction calling nextval() rolls back -- this is normal, not a bug.
GENERATED ... AS IDENTITY, covered next, manages a sequence for you automatically.

KEY TAKEAWAY A sequence generates guaranteed-unique, strictly increasing numbers, safely even under heavy concurrent access -- it's the engine behind every identity-column primary key.



PostgreSQL Sequences


A sequence is a database object that generates a unique, sequential number that can be used for primary keys or other unique identifiers.



Why Use Sequences?

- Generate unique values automatically
- Ensure no duplicate keys
- Simplify primary key generation
- Works well with high-concurrency applications

How Sequences Work



Common Sequence Operations

Operation	Example
Get next value	<code>SELECT nextval('employee_id_seq');</code>
Get current value	<code>SELECT currval('employee_id_seq');</code>
View sequence value	<code>SELECT last_value FROM employee_id_seq;</code>
Reset sequence	<code>ALTER SEQUENCE employee_id_seq RESTART WITH 1;</code>

Example Output

Operation	Result
nextval()	1
nextval()	2
nextval()	3
currval()	3
last_value	3

Creating and Using a Sequence

```
CREATE SEQUENCE employee_id_seq
START 1
INCREMENT 1
MINVALUE 1
CACHE 10;

CREATE TABLE employees (
  id BIGINT PRIMARY KEY DEFAULT nextval('employee_id_seq'),
  name VARCHAR(100),
  email VARCHAR(100)
);
```

Best Practices

- ✓ Use sequences for primary key generation.
- ✓ Set an appropriate increment value (typically 1).
- ✓ Use CACHE for better performance in high-load systems.
- ✓ Avoid manually updating sequence values.
- ✓ Monitor sequence usage in production.



Key Takeaway

PostgreSQL sequences provide a simple and reliable way to generate unique, sequential numbers, making them ideal for primary keys and other unique identifiers in your database schema.

IDENTITY COLUMNS

GENERATED AS IDENTITY is PostgreSQL's modern, SQL-standard way to auto-generate primary key values.

Form	Behavior
GENERATED ALWAYS AS IDENTITY	PostgreSQL always supplies the value -- an explicit INSERT value is rejected unless OVERRIDING SYSTEM VALUE is used.
GENERATED BY DEFAULT AS IDENTITY	PostgreSQL supplies the value only if the INSERT omits it -- an explicit value is allowed.


KEY TAKEAWAY GENERATED ALWAYS AS IDENTITY is the modern, standards-compliant default for PostgreSQL primary keys -- it manages a backing sequence for you automatically.



Identity Columns

Identity columns (**SQL standard**) automatically generate **unique values** for a column, typically used for primary keys. They are easier to use and manage than sequences + default.

How Identity Columns Work



Create Table with Identity Column

```
CREATE TABLE employees (
  employee_id BIGINT GENERATED ALWAYS AS IDENTITY
  PRIMARY KEY,
  first_name VARCHAR(50) NOT NULL,
  last_name VARCHAR(50) NOT NULL,
  email VARCHAR(100) UNIQUE,
  created_at TIMESTAMP DEFAULT now()
);
```

GENERATED ALWAYS
Database always
generates the value

AS IDENTITY
Marks column as
an identity column

Insert Without Providing ID

```
INSERT INTO employees (first_name, last_name, email)
VALUES
('John', 'Smith', 'john.smith@example.com');

INSERT INTO employees (first_name, last_name, email)
VALUES
('Sara', 'Johnson', 'sara.johnson@example.com');

INSERT INTO employees (first_name, last_name, email)
VALUES
('Michael', 'Brown', 'michael.brown@example.com');
```

Generated Values

employee_id (PK)	first_name	last_name	email
101	John	Smith	john.smith@example.com
102	Sara	Johnson	sara.johnson@example.com
103	Michael	Brown	michael.brown@example.com
...	...	...	...

Identity vs Sequence + Default

Identity Column	Sequence + Default
✓ SQL standard (since PostgreSQL 10) ✓ Simpler syntax and easier to use ✓ Automatically manages the underlying sequence ✓ Easier to alter (ADD/DROP IDENTITY) ✓ Preferred approach for new development	✓ Requires creating sequence separately ✓ More manual steps ✓ Default nextval('sequence_name') ✓ More flexible in some advanced scenarios ✓ Useful when sharing sequence across tables



Key Points

- Identity columns guarantee **unique, ever-increasing** values.
- Use **GENERATED ALWAYS** for values fully controlled by the database.
- Use **GENERATED BY DEFAULT** to allow explicit value insertion when needed.
- Identity columns are the **recommended** way to define auto-incrementing primary keys.



Best Practice: Use identity columns for primary keys unless you have a specific requirement that needs sequence sharing or advanced control.



PostgreSQL

CONNECTING TO POSTGRESQL WITH PSQL

psql is PostgreSQL's official command-line client -- the fastest way to run SQL directly against the server.

```
psql -h localhost -p 5432 -U crm_app -d crmdb
```

Meta-Command	Purpose
\l	List all databases on the server.
\c dbname	Connect to a different database.
\dt	List tables in the current schema's search path.
\d tablename	Describe a table -- columns, types, and constraints.
\q	Quit psql.

KEY TAKEAWAY psql is the fastest way to verify a schema exists and behaves as designed -- \dt and \d are the two meta-commands students will use most in the lab.



Connecting to PostgreSQL with `psql`

`psql` is the PostgreSQL command-line interface. It allows you to connect to a database server, run SQL queries, manage schemas, and perform administrative tasks.

1. Basic Connection Syntax

```
>_ psql -h <host> -p <port> -U <username> -d <database>
```



-h <host>

Database server host (default: localhost)



-p <port>

Port number (default: 5432)



-U <username>

Database user



-d <database>

Target database to connect to

2. Common Examples

```
>_ psql -U postgres # Connect to default 'postgres' database
```

```
>_ psql -U appuser -d empdb # Connect to empdb as appuser
```

```
>_ psql -h db.example.com -p 5432 -U appuser -d empdb # Remote connection
```

psql

```
$ psql -h localhost -p 5432 -U postgres -d empdb
```

Password for user postgres:

psql (16.2)

Type "help" for help.

```
empdb=# \conninfo
```

You are connected to database "empdb" as user "postgres" on host "localhost" (address "127.0.0.1") at port "5432".

```
empdb=# \dt
```

List of relations			
Schema	Name	Type	Owner
public	employee	table	postgres
public	department	table	postgres
public	job_role	table	postgres

(3 rows)

```
empdb=#
```

Useful psql Commands

<code>\l</code>	List all databases
<code>\c <db></code>	Connect to a database
<code>\dt</code>	List tables in current schema
<code>\d <table></code>	Describe table structure
<code>\du</code>	List database users
<code>\conninfo</code>	Show connection info
<code>\q</code>	Quit psql



Tip: Type `\?` to see all psql meta-commands.



Pro Tip

If you omit **-h**, **-p**, and **-d**, psql uses defaults:



Host: **localhost**



Port: **5432**



Database: **postgres**

PGADMIN OVERVIEW

pgAdmin is PostgreSQL's official graphical administration tool -- browse schemas, run queries, and inspect data visually.



Object Browser

Navigate servers, databases, schemas, and tables in a familiar tree view.



Query Tool

Write and run SQL with syntax highlighting and a results grid.



Table Inspector

View columns, constraints, and indexes without writing a single query.



EXPLAIN Visualizer

Renders a query's execution plan as a graphical diagram -- a Module 38 preview.

KEY TAKEAWAY pgAdmin gives a visual view of everything psql shows at the command line -- most PostgreSQL developers use both, switching based on the task.



pgAdmin Overview

pgAdmin is the official, open-source administration and development platform for PostgreSQL. It provides a powerful GUI to manage databases, objects, queries, users, permissions, and much more.

Key Capabilities

- Database Management**
Create, modify, and manage databases, schemas, tables, and other objects.
- Query Tool**
Write, run, and optimize SQL queries with syntax highlighting and results grid.
- Visual Tools**
ERD tool, query history, explain plans, and performance monitoring.
- Security & Access**
Manage roles, privileges, and object-level permissions.
- Backup & Restore**
Easily back up and restore databases and objects.



Tip: pgAdmin works on Windows, macOS, and Linux and connects to any PostgreSQL database.

pgAdmin Interface

The screenshot shows the pgAdmin interface with the following components:

- Menu & Toolbar:** Located at the top left, containing icons for various actions.
- Object Browser:** A tree view on the left showing the database structure (Servers, Databases, Schemas, Tables, Views, etc.).
- Dashboard Overview:** A central panel displaying various performance metrics and graphs, including Server sessions, Transactions per second, Tuples in/out, and Block I/O.
- Query Tool:** A panel for writing and executing SQL queries, showing a query history and a results grid.
- Results Grid:** A table displaying the results of the executed query, showing columns like employee_id, first_name, last_name, and department_name.

Common Tasks in pgAdmin

- Create Database**
Wizard-based database creation.
- Run SQL Queries**
Use the Query Tool to execute and analyze SQL.
- Browse & Edit Data**
View and edit table data directly.
- Manage Users & Roles**
Create roles and set permissions.
- Backup / Restore**
Export or restore databases with ease.
- Monitor Performance**
View activity, locks, and query statistics.



Why pgAdmin?

- Official PostgreSQL tool
- Feature-rich and extensible
- Ideal for DBAs, developers, and analysts



CONNECTING JAVA TO POSTGRESQL

A Java application talks to PostgreSQL through JDBC -- the same driver-based pattern used across every relational database.

```
String url = "jdbc:postgresql://localhost:5432/crmdb";
String user = "crm_app";
String pass = System.getenv("CRM_DB_PASSWORD");

try (Connection conn = DriverManager.getConnection(url, user, pass)) {
    // conn is a live PostgreSQL session
}
```


KEY TAKEAWAY Every JDBC connection needs three things: a jdbc:postgresql:// URL, credentials, and the PostgreSQL JDBC driver on the classpath -- Spring Boot automates all three from configuration.



Connecting Java to PostgreSQL

Java applications connect to PostgreSQL using **JDBC** (Java Database Connectivity). The PostgreSQL JDBC Driver provides the implementation for communicating between Java and the PostgreSQL database.

Connection Flow



3. Example: Connecting to PostgreSQL in Java

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.SQLException;

public class PostgresConnectionExample {
    public static void main(String[] args) {
        String url = "jdbc:postgresql://localhost:5432/employee_db";
        String user = "dbuser";
        String password = "securePassword";

        try (Connection conn = DriverManager.getConnection(url, user, password)) {
            System.out.println("🟢 Connected to PostgreSQL successfully!");
        } catch (SQLException e) {
            System.err.println("🔴 Connection failed: " + e.getMessage());
        }
    }
}
```



Tip: Use try-with-resources to automatically close connections and prevent resource leaks.

1. Add PostgreSQL JDBC Driver Dependency (Maven)

```
<dependency>
  <groupId>org.postgresql</groupId>
  <artifactId>postgresql</artifactId>
  <version>42.7.3</version>
</dependency>
```

2. JDBC Connection URL Structure

`jdbc:postgresql://<host>:<port>/<database>?<parameters>`

- `jdbc:postgresql` : JDBC protocol for PostgreSQL
- `host` : Database server hostname or IP
- `port` : Database port (default: 5432)
- `database` : Target database name
- `parameters` : Optional connection parameters

4. Common Connection Parameters

Parameter	Description	Example
ssl	Enable SSL connection	?ssl=true
connectTimeout	Connection timeout in seconds	?connectTimeout=10
currentSchema	Default schema to use	?currentSchema=public
ApplicationName	Name of the application	?ApplicationName=MyApp



Best Practices

- Use connection pooling (e.g., HikariCP) in production applications.
- Store database credentials in environment variables or configuration files.
- Always close connections, statements, and result sets.



PostgreSQL

POSTGRESQL JDBC DRIVER

The PostgreSQL JDBC driver is the library that translates Java's JDBC API calls into the PostgreSQL wire protocol.

```
<!-- Maven -->
<dependency>
  <groupId>org.postgresql</groupId>
  <artifactId>postgresql</artifactId>
  <scope>runtime</scope>
</dependency>
```

Registers itself with `java.sql.DriverManager` automatically -- no manual `Class.forName()` needed.
scope=runtime is intentional -- application code depends on the JDBC API, not this driver directly.
Spring Boot resolves the exact compatible version through its dependency management.


KEY TAKEAWAY Adding the `org.postgresql:postgresql` dependency is the one setup step every PostgreSQL Java project needs -- everything else in the connection flow is standard JDBC.



PostgreSQL JDBC Driver

The **PostgreSQL JDBC Driver** enables Java applications to connect to and interact with PostgreSQL databases.

Role in Java Applications



Add Dependency (Maven)

```
<dependency>
  <groupId>org.postgresql</groupId>
  <artifactId>postgresql</artifactId>
  <version>42.7.3</version>
</dependency>
```

Add Dependency (Gradle)

```
implementation 'org.postgresql:postgresql:42.7.3'
```

Example: Load the Driver and Get Connection

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.SQLException;

public class PostgreSQLConnectionExample {
    public static void main(String[] args) {
        String url = "jdbc:postgresql://localhost:5432/employeesdb";
        String user = "dbuser";
        String password = "securePassword";

        try (Connection conn = DriverManager.getConnection(url, user, password)) {
            System.out.println("🟢 Connected to PostgreSQL successfully!");
        } catch (SQLException e) {
            System.err.println("🔴 Connection failed: " + e.getMessage());
        }
    }
}
```



Tip: Use the official PostgreSQL JDBC Driver for maximum compatibility and performance.

JDBC Driver Types

Driver Type	Description	PostgreSQL Support
Type 1	JDBC-ODBC Bridge Driver	Not supported
Type 2	Native-API (partially Java)	Not supported
Type 3	Network Protocol Driver	Not supported
Type 4	Pure Java Driver (Thin Driver)	✔ Supported (Recommended)

Why Type 4 (Pure Java) Driver?

- ✔ 100% Java – platform independent
- ✔ High performance and lightweight
- ✔ No native libraries required
- ✔ Fully supports PostgreSQL features
- ✔ Recommended and actively maintained

Common Driver Class

For PostgreSQL Type 4 driver, the class is automatically registered via SPI.

```
org.postgresql.Driver
```

Driver JAR

The PostgreSQL JDBC Driver is packaged as a JAR file.



postgresql-42.7.3.jar
(Driver JAR)

Included in



Application
Classpath



Best Practices

- Always use the latest stable driver version.
- Store the JAR in your local repository or dependency manager.
- Use try-with-resources to ensure connections are closed.
- Do not hardcode credentials; use configuration files or environment variables.



PostgreSQL

CONNECTION URL STRUCTURE

Every PostgreSQL JDBC URL follows the same four-part shape: protocol, host, port, and database.

Part	Example	Meaning
Protocol	jdbc:postgresql://	Tells DriverManager which driver should handle this URL.
Host	localhost	Where the PostgreSQL server is running.
Port	5432	PostgreSQL's conventional port -- can be omitted if using the default.
Database	/crmdb	Which database on that server to connect to.

KEY TAKEAWAY `jdbc:postgresql://host:port/database` -- get this exact shape memorized; Module 39's `application.properties` `datasource` URL uses this identical format.




Connection URL Structure

A JDBC connection URL tells the PostgreSQL JDBC Driver how to connect to your database.

General Format

`jdbc:postgresql://host:port/database??param1=value1¶m2=value2`



Common Connection Parameters

Parameter	Description	Example
<code>ssl</code>	Enable SSL encryption	<code>?ssl=true</code>
<code>connectTimeout</code>	Connection timeout in seconds	<code>?connectTimeout=10</code>
<code>socketTimeout</code>	Socket read timeout in seconds	<code>?socketTimeout=30</code>
<code>ApplicationName</code>	Name of the application	<code>?ApplicationName=MyApp</code>
<code>currentSchema</code>	Default schema to use	<code>?currentSchema=public</code>
<code>loginTimeout</code>	Login timeout in seconds	<code>?loginTimeout=5</code>



Best Practices

- Use environment variables or configuration files for connection URLs.
- Always enable SSL (`ssl=true`) in production environments.
- Set appropriate timeouts to avoid hanging connections.
- Use `ApplicationName` to identify your application in PostgreSQL logs.



Quick Tip

If no port is specified, the driver uses the default PostgreSQL port **5432**. Keep your credentials secure—never hardcode them in source code.

Example Connection URLs

Description	JDBC Connection URL
Local default connection	<code>jdbc:postgresql://localhost:5432/mydb</code>
Local with custom port	<code>jdbc:postgresql://localhost:6543/mydb</code>
Remote host	<code>jdbc:postgresql://dbserver.example.com:5432/mydb</code>
With parameters	<code>jdbc:postgresql://localhost:5432/mydb?ssl=true&connectTimeout=10&ApplicationName=MyApp</code>
Using IP address	<code>jdbc:postgresql://192.168.1.50:5432/mydb</code>

Where Each Part Comes From



Host & Port

Provided by your DBA or database administrator.



Database

The name of the database you want to connect to.



Parameters

Optional settings to control connection behavior.



PostgreSQL

DATABASE USERS AND ROLES

PostgreSQL has a single unified concept, the role, that represents both a user and a group.

```
CREATE ROLE crm_app WITH LOGIN PASSWORD 'change_me' NOSUPERUSER;  
  
-- roles can be grouped  
CREATE ROLE crm_readonly;  
GRANT crm_readonly TO crm_app;
```

LOGIN makes a role usable for connecting -- roles without LOGIN act as pure groups.
A superuser role bypasses every permission check -- never used by an application.
\du in psql lists every role on the server.

KEY TAKEAWAY PostgreSQL calls both individual users and groups 'roles' -- CREATE ROLE ... WITH LOGIN is how an application, not a human, gets its own database identity.



Database Users and Roles

Control access to PostgreSQL resources using roles. Users are roles that can log in.



Key Concepts

- **Role:** A named entity that can own objects and have privileges.
- **User:** A role that can authenticate (has LOGIN privilege).
- **Privileges:** Permissions to perform actions on database objects.
- **Ownership:** Object owners automatically have all privileges.
- **Principle of Least Privilege:** Grant only the permissions needed.




Common Privileges

Privilege	Description
CONNECT	Allows connecting to a database
USAGE	Allows using a schema
SELECT	Allows reading data from tables/views
INSERT	Allows inserting data
UPDATE	Allows updating data
DELETE	Allows deleting data
ALL PRIVILEGES	Grants all privileges on an object



Role Hierarchy and Access



Example SQL

Create Roles

```
CREATE ROLE app_user LOGIN PASSWORD 'App@123';
CREATE ROLE readonly_user LOGIN PASSWORD 'Read@123';
CREATE ROLE reporting_user LOGIN PASSWORD 'Report@123';
CREATE ROLE app_role;

GRANT app_role TO app_user;
```

Grant Privileges

```
GRANT CONNECT ON DATABASE employee_db TO app_role;
GRANT USAGE ON SCHEMA public TO app_role;
GRANT SELECT, INSERT, UPDATE, DELETE
ON ALL TABLES IN SCHEMA public TO app_role;
GRANT SELECT ON ALL TABLES IN SCHEMA public
TO readonly_user, reporting_user;
```



Best Practices



Follow Least Privilege
Grant only necessary permissions.



Use Roles for Groups
Assign users to roles instead of granting privileges directly.



Avoid Superuser for Applications
Use dedicated roles.



Review Privileges Regularly
Audit access and remove unnecessary permissions.



Separate Duties
Use different roles for development, reporting, and administration.

PERMISSIONS AND LEAST PRIVILEGE

Grant a role exactly the permissions it needs to do its job, and nothing more.

```
GRANT CONNECT ON DATABASE crmdb TO crm_app;  
GRANT USAGE ON SCHEMA public TO crm_app;  
GRANT SELECT, INSERT, UPDATE ON customer, account TO crm_app;  
GRANT USAGE ON ALL SEQUENCES IN SCHEMA public TO crm_app;
```

KEY TAKEAWAY An application role gets exactly the GRANTs its code path needs -- CONNECT, USAGE, and table-level SELECT/INSERT/UPDATE, never superuser or DROP/DELETE it doesn't use.

TRY IT YOURSELF — EXERCISE 5: PLAN THE ROLE AND CONNECTION

Reinforces: psql, the JDBC driver, connection URL structure, and the users, roles, and least-privilege slides you just covered.

STEPS (notes/lab37-role-connection.md)

Step	What To Do
1 — JDBC URL	Write the connection URL shape for the CRM database, with no password in it.
2 — App Role	Say what crm_app may do and what it must not.
3 — Least Privilege	Give one thing the app role is denied and the reason.
4 — Secret Handling	Say where the password lives, given it is never in Git.

Role and connection

```
jdbc:postgresql://localhost:5432/crm user=crm_app (password injected)
grant select, insert, update, delete on crm tables to crm_app
no ddl, no superuser, no drop -- migrations run as a different role
```

TRY IT NOW — Complete Exercise 5 before continuing: exercises/exercise-05-role-and-connection.md (workspace: examples/module-37-exercises).



Permissions and Least Privilege

Grant only the permissions required. **Nothing more.**



What Are Permissions?

Permissions (privileges) **control what actions** a user or role can perform on database objects.



Permission Levels

- Object Level**
Apply to specific objects (tables, views, sequences, functions).
- Schema Level**
Apply to all objects within a schema.
- Database Level**
Apply to the entire database.
- System Level**
Apply to database-wide administrative actions.



Common Object Privileges

Privilege	Description
SELECT	Read data from tables or views
INSERT	Insert new rows
UPDATE	Modify existing rows
DELETE	Remove rows
REFERENCES	Create foreign keys
EXECUTE	Execute functions or procedures
USAGE	Use schemas, sequences, types



Example: Applying Least Privilege

Application role needs to read and write employee data, but should not be able to drop or alter objects.

```
CREATE ROLE app_user LOGIN PASSWORD 'App@123';
```

Grant	Purpose
GRANT USAGE ON SCHEMA public TO app_user;	Allows use of the public schema
GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE public.employees TO app_user;	Allows CRUD on employees (table only)
GRANT SELECT ON TABLE public.departments TO app_user;	Allows read-only access to departments
GRANT USAGE, SELECT ON SEQUENCE public.employee_id_seq TO app_user;	Allows use of the sequence for IDs
GRANT EXECUTE ON FUNCTION public.calculate_bonus() TO app_user;	Allows executing a specific function



Result

The app_user role can **only** perform what it is explicitly granted. It cannot alter, drop, or access objects it does not have permission to use.



Least Privilege in Action

Least Privilege (Good) ✓

app_user
(Application Role)

- ✓ USAGE on schema
- ✓ SELECT, INSERT, UPDATE, DELETE on needed tables
- ✓ USAGE on required sequences
- ✓ EXECUTE on specific functions



Access Only to Required Objects and Actions
Reduced risk and impact

Over-Privileged (Risky) ✗

app_user
(Over-Privileged)

- ✗ ALL PRIVILEGES on schema
- ✗ ALL PRIVILEGES on all tables
- ✗ ALL PRIVILEGES on sequences
- ✗ CREATE, DROP, ALTER, ...



Access Beyond Requirements
Higher risk and potential damage



Key Takeaway

Grant the **minimum permissions** needed for the job. Review permissions regularly and remove what is no longer required.



Best Practices



Use Roles, Not Users
Assign permissions to roles and grant roles to users.



Apply Least Privilege
Grant only what is necessary to perform the task.



Review Regularly
Audit permissions and remove unnecessary access.



Monitor and Log
Enable logging to track access and changes.



Separate Duties
Use different roles for development, reporting, and administration.

TRANSACTIONS OVERVIEW

A transaction groups one or more statements into a single all-or-nothing unit of work.

```
BEGIN;  
UPDATE account SET balance = balance - 100 WHERE account_id = 1;  
UPDATE account SET balance = balance + 100 WHERE account_id = 2;  
COMMIT;  -- or ROLLBACK; to undo both statements
```

KEY TAKEAWAY BEGIN...COMMIT groups statements so they succeed or fail together -- essential the moment more than one row must change consistently, like a transfer between two accounts.



Transactions Overview

Transactions ensure **data integrity** by treating a series of operations as a single unit of work.



What is a Transaction?

A transaction is a sequence of one or more SQL operations executed as a single logical unit of work. Either all operations succeed, or none are applied.



Key Properties

- Maintains data consistency and reliability
- Provides isolation between concurrent operations
- Ensures database remains in a valid state
- Controlled with **COMMIT**, **ROLLBACK**, and **SAVEPOINT**




Transaction Lifecycle



If any operation fails, use **ROLLBACK** to prevent partial updates and keep the database consistent.



How a Transaction Works



Example (PostgreSQL)

Successful Transaction

```
BEGIN;  
UPDATE accounts  
  SET balance = balance - 500  
  WHERE account_id = 1001;  
  
UPDATE accounts  
  SET balance = balance + 500  
  WHERE account_id = 1002;  
  
COMMIT; -- Changes are saved
```



Failed Transaction (Rolled Back)

```
BEGIN;  
UPDATE accounts  
  SET balance = balance - 500  
  WHERE account_id = 1001;  
  
UPDATE accounts  
  SET balance = balance + 500  
  WHERE account_id = 9999; -- Invalid account  
  
ROLLBACK; -- No changes applied
```



ACID Properties



Atomicity

All operations in the transaction succeed or none do.



Consistency

Database remains in a valid state before and after the transaction.



Isolation

Concurrent transactions do not interfere with each other.



Durability

Once committed, changes are permanent, even after failures.



Common Use Cases

- Money transfers between accounts
- Order processing
- Inventory updates
- Bulk data operations
- Any operation requiring data consistency



Best Practices



Keep transactions short
Minimize lock contention and improve performance.



Do not mix user interaction within transactions
Collect input first, then start the transaction.



Handle exceptions properly
Rollback on error to maintain consistency.



Use appropriate isolation level
Balance consistency needs with performance.



Monitor long-running transactions
They can block other operations.

ACID IN POSTGRESQL

ACID is the four-part guarantee PostgreSQL makes about every transaction: Atomicity, Consistency, Isolation, Durability.



Atomicity

All of a transaction's statements succeed together, or none of them do -- no partial results.



Consistency

A transaction can only move the database from one valid state to another -- constraints are never violated.



Isolation

Concurrent transactions don't see each other's uncommitted changes -- enforced by MVCC.



Durability

Once COMMIT returns, the change survives a crash -- it's written to the WAL before confirming.

KEY TAKEAWAY PostgreSQL's ACID guarantees are what let a developer trust that a committed transaction really happened, and a rolled-back one really didn't -- even after a crash.



ACID in PostgreSQL

PostgreSQL ensures reliable transactions using **ACID** properties.



What is ACID?

ACID is a set of properties that guarantee reliable processing of database transactions, even in the event of errors, failures, or concurrent access.



PostgreSQL is fully ACID-compliant.

It uses MVCC, Write-Ahead Logging (WAL), and robust locking mechanisms to protect your data.



Why ACID Matters

- Prevents data corruption
- Maintains consistency under concurrent load
- Supports safe recovery from failures
- Builds trust in your data and applications



The ACID Properties

A

Atomicity

All operations in a transaction are treated as a single unit. Either all succeed or none do.

**C**

Consistency

A transaction brings the database from one valid state to another, preserving all rules and constraints.

**I**

Isolation

Concurrent transactions do not interfere with each other. Each transaction appears to run in isolation.

**D**

Durability


Once a transaction is committed, changes are permanent—even if the system crashes.



How PostgreSQL Achieves ACID

- MVCC (Multi-Version Concurrency Control) for isolation
- WAL (Write-Ahead Logging) for durability
- Locks and row-level locking for consistency and isolation
- Crash recovery and checkpoints for atomicity and durability

Example: Bank Transfer Transaction



Example in PostgreSQL (SQL)

```
BEGIN;
UPDATE accounts
  SET balance = balance - 200
  WHERE id = 1; -- Debit Account A

UPDATE accounts
  SET balance = balance + 200
  WHERE id = 2; -- Credit Account B
```

```
-- If everything is correct
COMMIT; -- Changes are saved

-- If any error occurs
ROLLBACK; -- No changes applied
```



Best Practices



Keep Transactions Short

Long transactions increase lock contention and impact performance.



Do Not Mix User Interaction

Collect input first, then start the transaction.



Handle Exceptions

Use ROLLBACK on errors to maintain consistency.



Test Failures

Regularly test rollback and recovery scenarios.



Monitor Performance

Track long-running transactions and lock waits.

ENTERPRISE SCHEMA DESIGN BEST PRACTICES

Six practices that separate a schema that merely works from one that survives production.

Practice	In Short
Business Requirements	Model real entities and rules -- design for today and for growth.
Model Data Correctly	Right types, normalized to 3NF, surrogate keys.
Data Integrity	PK, FK, NOT NULL, UNIQUE, CHECK enforced at the database level.
Index Smartly	Index foreign keys and hot filter columns -- avoid over-indexing.
Scalability & Security	Plan for growth; enforce least-privilege roles.
Maintainability & Governance	Consistent naming, documented schema, clear ownership.

KEY TAKEAWAY Every enterprise schema decision in this course traces back to one of these six practices -- treat this table as a checklist to review before any production migration.



Enterprise Schema Design Best Practices

Build **scalable, maintainable, secure, and high-performance** PostgreSQL schemas.



Guiding Principles



Model the Business

Design schemas that reflect business entities, rules, and processes.



Plan for Scale

Design for growth in data volume, users, and performance.



Ensure Data Integrity

Use constraints and relationships to enforce correctness.



Keep it Simple

Avoid unnecessary complexity and premature optimization.



Secure by Design

Apply least privilege and protect sensitive data.



Good Schema Design Benefits

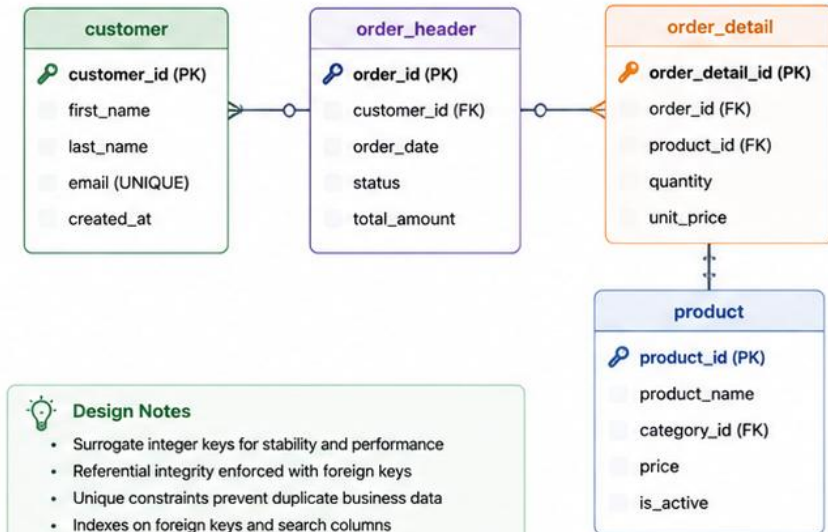
- ✓ Consistent and accurate data
- ✓ High performance and scalability
- ✓ Easier maintenance and evolution
- ✓ Stronger security and compliance
- ✓ Lower total cost of ownership



Best Practices Checklist

- ✓ Use **meaningful**, consistent naming for schemas, tables, and columns.
- ✓ Follow **normalization** (3NF) to reduce redundancy and anomalies.
- ✓ Define **primary keys** for all tables using surrogate keys when appropriate.
- ✓ Use **foreign keys** to enforce relationships and maintain referential integrity.
- ✓ Apply appropriate **data types** and sizes—avoid oversized columns.
- ✓ Enforce NOT NULL, UNIQUE, CHECK, and DEFAULT constraints.
- ✓ Create indexes strategically on joins, filters, and unique columns.
- ✓ **Avoid over-indexing**—review and remove unused indexes.
- ✓ Design for **auditability** (created_by, created_at, updated_at, etc.).
- ✓ Plan for **data lifecycle**—archiving, purging, and retention policies.
- ✓ Use **schemas** to organize domains and control access.
- ✓ **Document** the data model, rules, and relationships.

Example: Well-Designed Schema



Design Notes

- Surrogate integer keys for stability and performance
- Referential integrity enforced with foreign keys
- Unique constraints prevent duplicate business data
- Indexes on foreign keys and search columns
- Clear separation of customer, order, and product domains



Common Pitfalls to Avoid

- ✗ Using reserved words or unclear names
- ✗ Storing multiple values in a single column
- ✗ Skipping constraints or primary keys
- ✗ Over-normalizing creating unnecessary complexity
- ✗ Missing or poorly designed indexes
- ✗ Ignoring growth, archiving, and retention needs



Additional Recommendations



Review and refactor schemas regularly to adapt to changing business needs.



Use naming standards (e.g., snake_case) and consistent conventions across the organization.



Add comments to tables and columns to explain purpose and business meaning.



Monitor query performance and adjust schema, indexes, and statistics as needed.



Leverage tools like pgAdmin, ER diagrams, and migration frameworks (Flyway/Liquibase).



Collaborate with stakeholders, DBAs, and developers during design and reviews.

COMMON DATABASE DESIGN MISTAKES

The same handful of mistakes account for most schema problems found in real code review.

Mistake	Why It Hurts	Fix
Storing money as REAL/FLOAT	Rounding error corrupts balances over time.	Use NUMERIC(19,2).
No foreign keys ('app enforces it')	Orphan rows appear the moment any code path has a bug.	Declare real FOREIGN KEY constraints.
Everything nullable by default	Hides which data is truly required.	Add NOT NULL wherever a business rule requires it.
Indexing every column	Slows every write; wastes storage.	Index with intent, based on real query patterns.
No naming convention	Slows every future change and review.	Consistent snake_case, documented schema.

KEY TAKEAWAY Each row in this table is a real mistake this module already showed how to avoid -- treat it as a pre-submission checklist for Lab 37.

TRY IT YOURSELF — EXERCISE 6: LAB 37 READINESS SELF-CHECK

Pre-lab gate: sweep your schema against the common-mistakes slide, plan the negative tests, and confirm the five earlier notes files are complete.

STEPS (notes/lab37-readiness.md)

Step	What To Do
1 — Mistake Sweep	Check your schema against the common-mistakes slide and fix what matches.
2 — Negative Tests	Name three inserts that must fail, and the constraint that rejects each.
3 — Rebuild Proof	Say how you will prove the schema can be dropped and recreated cleanly.
4 — Pass Mark	Mark Pass only if all five earlier notes files are complete, not stubbed.

Negative tests

```
duplicate email      -> unique violation
status 'PENDING'    -> check violation
account with unknown customer_id -> foreign key violation
```

TRY IT NOW — Complete Exercise 6 before Lab 37: exercises/exercise-06-lab37-readiness.md (workspace: examples/module-37-exercises).



Common Database Design Mistakes

Avoid these pitfalls to build **scalable, reliable, and high-performance** PostgreSQL databases.

Mistake	Why It's a Problem	Better Approach
 1. Poor or Missing Primary Keys Tables without primary keys or with non-unique keys.	✗ Leads to duplicate data, update anomalies, poor performance, and replication issues.	✓ Always define a primary key using surrogate keys (e.g., BIGSERIAL/UUID) when appropriate.
 2. Over-Normalization Splitting data into too many tables without business need.	✗ Increases complexity, more joins, and slower queries.	✓ Normalize to 3NF, but denormalize strategically for performance when justified.
 3. Using SELECT * Everywhere Retrieving unnecessary columns in queries.	✗ Consumes more I/O and network bandwidth. Impacts query performance.	✓ Select only the required columns. Be explicit in queries and APIs.
 4. Incorrect Data Types Using generic or oversized data types (e.g., TEXT for IDs).	✗ Wastes storage, slows down queries, and reduces data integrity.	✓ Use appropriate data types (e.g., INTEGER, TIMESTAMP, NUMERIC, VARCHAR(n)).
 5. Missing or Poorly Designed Indexes No indexes on key columns or wrong index strategy.	✗ Causes full table scans and poor performance.	✓ Index foreign keys and frequently filtered, sorted, or joined columns. Avoid over-indexing.
 6. NULL Overuse Allowing NULLs when not necessary.	✗ Complicates logic, causes unexpected results, and wastes space.	✓ Use NOT NULL with appropriate DEFAULT values where possible.
 7. Weak Constraints Missing constraints (FK, UNIQUE, CHECK) on tables.	✗ Allows invalid data and breaks referential integrity.	✓ Define FOREIGN KEY, UNIQUE, CHECK, and DEFAULT constraints consistently.
 8. Ignoring Query Performance Early Designing schema without considering query patterns.	✗ Leads to redesign, high query latency, and scalability issues.	✓ Understand query patterns early. Design schema, indexes, and partitions accordingly.
 9. No Archiving or Partitioning Plan Storing all data in a single table indefinitely.	✗ Leads to huge tables, slow queries, and maintenance challenges.	✓ Plan for partitioning, archiving, and retention policies from the start.
 10. Mixing Business Rules in the Database Embedding complex business logic in schema (triggers/procedures).	✗ Hard to maintain, test, and evolve. Tight coupling.	✓ Keep business logic in the application layer. Use the database for data integrity only.



The Impact of Poor Design



Poor Performance and Scalability



Data Inconsistency and Integrity Issues



Higher Development and Maintenance Cost



Difficult Reporting and Analytics



Harder to Evolve and Extend



Good Design Outcomes



High Performance and Scalability



Accurate and Consistent Data



Lower Development and Maintenance Cost



Easier Reporting and Analytics



Future-Proof and Flexible Schema



Key Takeaway

Good database design is the foundation of a reliable and high-performing application. Plan well, validate early, and review continuously.



Best Practices Reminder



Understand your business and data before designing.



Keep it simple, consistent, and well-documented.



Review and test queries during design.



Monitor performance and adjust indexes and schema as needed.



Collaborate with DBAs, architects, and developers during design reviews.

LAB OVERVIEW — POSTGRESQL DATABASE FUNDAMENTALS

Create CRM DB/schema/tables with PKs, FKs, constraints, indexes, seeds /

BUSINESS SCENARIO

- Angular (Labs 33–36) talks to Spring REST; Spring will persist to PostgreSQL. Leadership freezes:

LEARNING OBJECTIVES

- Create a PostgreSQL database/schema and CRM tables with named constraints Model one-to-many Customer→Account with FK indexes Seed deterministic fixtures (`CUS-1001`, `CUS-1002`) and prove negatives Connect with `psql`/pgAdmin and document a Java JDBC URL Apply least-privilege roles and state ACID properties for CRM writes

WHAT YOU BUILD

- `examples/lab37-crm/` — SQL scripts, compose, design notes | Named constraints · CUS-1001/CUS-1002 seeds · one negative check · notes |

KEY TAKEAWAY Design and implement the ****PostgreSQL** CRM schema before ORM: databases vs schemas, tables, PKs/FKs, UNIQUE/NOT NULL/CHECK, indexes, roles/least privilege, ACID transactions, `psql`/pgAdmin, and the JDBC URL shape Spring Boot will use in Lab 39.**



Lab Overview – PostgreSQL Database Design

Design a **normalized**, **secure**, and **performant** schema in PostgreSQL for a real-world business scenario.



Business Scenario

You are designing the database for an **Online Learning Platform** where students can enroll in courses, instructors create content, and payments are processed.

- ✓ Students can register and enroll in courses
- ✓ Instructors create courses and lessons
- ✓ Students make payments for paid courses
- ✓ Admins manage users and content

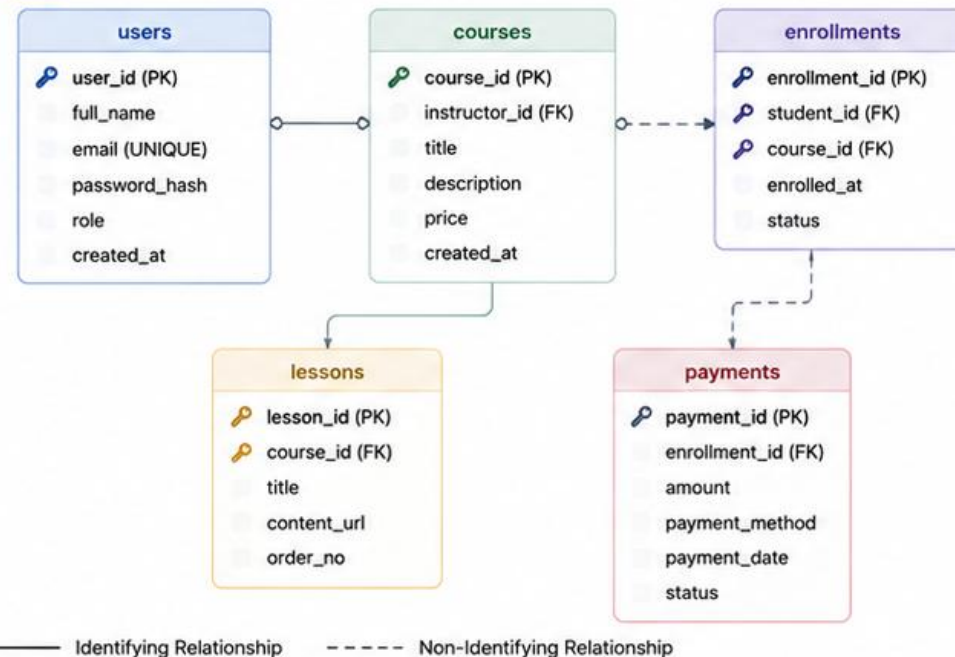


Lab Objectives

- ✓ Design a normalized schema (3NF)
- ✓ Define tables with appropriate data types
- ✓ Implement primary keys, foreign keys, and constraints
- ✓ Create indexes for performance
- ✓ Apply best practices for security and maintainability
- ✓ Populate the database with sample data
- ✓ Write queries to validate design and performance



High-Level Schema



Lab Environment



Database
PostgreSQL 15+



Tools
pgAdmin / psql



Provided
Sample data sets, SQL scripts, and requirements document



Deliverables
SQL script, ER diagram, queries, and report



What You Will Practice

- Creating schemas and tables
- Applying constraints and relationships
- Choosing proper data types
- Indexing strategy
- Normalizing data
- Writing queries and validating performance



Lab Tasks

1

Requirements Review

Understand business requirements and identify entities.



2

Conceptual Design

Create ER diagram and identify relationships and cardinalities.



3

Logical Design

Convert ERD to relational schema and normalize to 3NF.



4

Physical Design

Define data types, constraints, indexes, and sequences.



5

Implementation

Create schema in PostgreSQL using SQL scripts.



6

Validation & Testing

Insert sample data, write queries, and analyze performance.



7

Documentation

Prepare ERD, data dictionary, and design rationale.



Key Learning Outcomes



Design normalized schemas that reduce redundancy.



Implement constraints that ensure data integrity.



Create indexes that improve query performance.



Apply security best practices with roles and least privilege.



Evaluate and optimize queries using EXPLAIN ANALYZE.

LAB TASKS AND DELIVERABLES

PostgreSQL Database Fundamentals — implementation steps, deliverables, and the evaluation rubric.

#	Implementation Step
1	Open starter/README.md.
2	Copy starter/ into java-bootcamp/examples/lab37-crm.
3	Fill every -- TODO in database/*.sql.
4	Smoke: compose up → apply scripts → seed SELECT → one negative check.
5	Evidence under notes/screenshots/lab-37/.
6	ER notes with cardinalities (database/er-diagram.md)
7	PostgreSQL Docker (or shared) runtime
8	Least-privilege crm_app role script
9	DDL: CUSTOMER + ACCOUNT (+ indexes); full path: ADDRESS + HISTORY
10	Seed Amina CUS-1001 ACTIVE + account; Ravi CUS-1002 PROSPECT
11	Negative verify (duplicate email / bad status / orphan FK)
12	Drop/recreate proof + design-decisions.md

EXPECTED DELIVERABLES

- ER notes with cardinalities (database/er-diagram.md) PostgreSQL Docker (or shared) runtime Least-privilege crm_app role script DDL: CUSTOMER + ACCOUNT (+ indexes); full path: ADDRESS + HISTORY Seed Amina CUS-1001 ACTIVE + account; Ravi CUS-1002 PROSPECT Negative verify (duplicate email / bad status / orphan FK) Drop/recreate proof + design-decisions.md JDBC URL note for Java; no passwords in Git

RUBRIC (100 MARKS)

- Environment 10 · Core Impl 30 · Integration/Config 15 · Failure Handling 15 · Verification 10 · Security/Prod Awareness 10 · Docs/Evidence 10

KEY TAKEAWAY Fixtures CUS-1001 (Amina Khan, ACTIVE) and CUS-1002 (Ravi Singh, PROSPECT) anchor Lab 37; complete pre-lab exercises 1→6 first, then the graded lab.



Lab Tasks and Deliverables

Complete the tasks to **design, implement, and validate** a robust PostgreSQL database schema.



Lab Tasks

- 1 Requirements Review**
Review the business scenario and identify key entities, attributes, and rules.
- 2 Conceptual Design (ERD)**
Create an ER diagram showing entities, relationships, and cardinalities.
- 3 Logical Design**
Map the ERD to a relational schema and normalize to 3NF.
- 4 Physical Design**
Define tables with appropriate data types, primary keys, foreign keys, constraints, indexes, and sequences.
- 5 Schema Implementation**
Write and execute SQL scripts to create schemas, tables, constraints, indexes, and sequences.
- 6 Data Population**
Insert provided sample data and any additional lookup/reference data.
- 7 Validation & Testing**
Run queries to validate relationships, constraints, and performance.
- 8 Performance Analysis**
Analyze slow queries using EXPLAIN ANALYZE and add indexes if needed.
- 9 Documentation**
Document the schema, data dictionary, and design decisions.



Deliverables

Deliverable	Descriptions
ER Diagram	High-quality ER diagram (PNG/PDF) showing entities, relationships, and cardinalities.
SQL Scripts	Complete SQL scripts to create schemas, tables, constraints, indexes, sequences, and sample data.
Data Dictionary	Data dictionary listing tables, columns, data types, constraints, keys, and descriptions.
Sample Queries	SQL queries demonstrating CRUD operations, joins, and business rules.
Performance Report	Query performance analysis using EXPLAIN ANALYZE with observations and improvements.
Design Rationale	Document summarizing design decisions, assumptions, and trade-offs.
README	Instructions to set up and run the database and scripts.



Success Criteria

- ✓ Schema is normalized to 3NF with minimal redundancy.
- ✓ All constraints and relationships work as expected.
- ✓ Sample data inserts without errors.
- ✓ Queries return correct results and perform efficiently.
- ✓ Documentation is clear, complete, and well-structured.



Evaluation Focus

- Correctness**
Accurate model, constraints, and relationships.
- Normalization**
Proper normalization to 3NF.
- Data Integrity**
Appropriate constraints and referential integrity.
- Performance**
Effective indexing and query performance.
- Documentation**
Quality and completeness of documentation.
- Best Practices**
Use of naming standards, data types, and design principles.



Helpful Tips

- Start with a clear understanding of the business rules.
- Draw the ERD before writing any SQL.
- Choose the right data types for accuracy and performance.
- Define keys and constraints early.
- Use meaningful names for schemas, tables, and columns.
- Validate with real queries and edge cases.



Recommended Workflow



MODULE 37 SUMMARY

You can now design, constrain, and connect to a normalized PostgreSQL schema.

KEY CONCEPTS COVERED



Fundamentals & Architecture

Relational concepts and how PostgreSQL organizes databases, schemas, and processes.



Data Types

Numeric, character, boolean, date/time, UUID, and JSON/JSONB types.



Keys & Constraints

Primary keys, foreign keys, UNIQUE, NOT NULL, CHECK, and DEFAULT values.



Relationships & Normalization

One-to-one, one-to-many, many-to-many, and reducing redundancy.



Indexes & Sequences

B-Tree, unique, and composite indexes; sequences and identity columns.



Connectivity & Governance

psql, pgAdmin, JDBC, users, roles, transactions, and ACID.

MODULE FLOW RECAP

1

1. Model

2

2. Constrain

3

3. Relate

4

4. Index

5

5. Connect

KEY TAKEAWAY You can now design a normalized PostgreSQL schema, enforce its rules with constraints, index it for performance, and connect to it from psql and Java.



MODULE SUMMARY

Module 13: REST API Design with Java



In this module, you explored the principles and best practices for designing high-quality, scalable, and consistent RESTful APIs using Java and HTTP. You learned how to model resources, design clean URIs, choose the right HTTP methods, structure requests and responses, handle errors, and document APIs using OpenAPI.

WHAT YOU LEARNED



REST Principles
and Constraints




Client-Server
and Stateless
Communication



Resource-Based
API Design



HTTP Methods
and Status Codes



Request & Response
JSON and Headers



Pagination,
Filtering, and
Sorting



Error Handling
and API Error
Design



API Versioning
Strategies



OpenAPI
and API-First
Development



Best Practices
and Common
Mistakes



KEY TAKEAWAYS



REST APIs enable scalable, interoperable, and maintainable enterprise integrations.



A well-designed API starts with clear resources, clean URIs, and correct use of HTTP methods.



Consistent request/response structures and meaningful status codes improve reliability and developer experience.



Pagination, filtering, sorting, and versioning are essential for real-world APIs.



OpenAPI documentation ensures clarity, standardization, and collaboration.



Following best practices and avoiding common mistakes leads to robust and sustainable APIs.



You are now equipped with the foundational knowledge to design RESTful APIs that are consistent, scalable, secure, and aligned with enterprise standards.

KNOWLEDGE CHECK (1 OF 2)

Test your understanding of relational fundamentals, data types, and keys.

1. Which PostgreSQL object groups related tables under a shared namespace inside one database?

- A. A tablespace
- B. A schema**
- C. A sequence
- D. A role

2. Which PostgreSQL data type should store exact monetary amounts?

- A. REAL
- B. FLOAT
- C. NUMERIC(p,s)**
- D. TEXT

3. What must be true of every primary key value in a table?

- A. It may repeat across rows
- B. It may be NULL for optional rows
- C. It must be unique and NOT NULL**
- D. It must be a UUID

4. A foreign key column in a child table must:

- A. Always be NOT NULL
- B. Reference an existing primary key value in the parent table (or be NULL)**
- C. Be indexed automatically by PostgreSQL
- D. Match the child table's own primary key

KEY TAKEAWAY Schemas group tables inside a database; NUMERIC stores exact money; a primary key is always unique and NOT NULL; a foreign key must match an existing parent key or be NULL.

KNOWLEDGE CHECK (2 OF 2)

Four more questions on relationships, indexes, transactions, and connecting to PostgreSQL.

1. A many-to-many relationship between ORDER and PRODUCT is correctly modeled in PostgreSQL by:

- A. Adding a product_id column directly to ORDER
- B. A junction table with foreign keys to both tables**
- C. Storing a comma-separated list of product IDs
- D. Using a CHECK constraint

3. What does the ACID guarantee 'Atomicity' mean in PostgreSQL?

- A. Queries run in alphabetical order
- B. A transaction's statements all succeed or all roll back together**
- C. Every column value is atomic and indivisible
- D. Only one user can connect at a time

2. Which index type does PostgreSQL use by default for CREATE INDEX?

- A. Hash
- B. GIN
- C. B-Tree**
- D. BRIN

4. Which JDBC connection URL correctly targets a PostgreSQL database named crmdb on localhost?

- A. jdbc:mysql://localhost:5432/crmdb
- B. postgresql://localhost/crmdb
- C. jdbc:postgresql://localhost:5432/crmdb**
- D. jdbc:postgres:crmdb@localhost

KEY TAKEAWAY Many-to-many relationships need a junction table; B-Tree is PostgreSQL's default index; Atomicity means a transaction commits or rolls back as one unit; the JDBC URL form is jdbc:postgresql://host:port/database.



KNOWLEDGE CHECK

Module 13: REST API Design with Java



Test your understanding of key concepts from this module.

1 Which REST constraint means the server does not store any client context between requests?

- ☐ A Cacheability
- ☐ B Client-Server
- ☐ C Stateless
- ☐ D Uniform Interface



2 Which HTTP method should be used to create a new resource?

- ☐ A GET
- ☐ B POST
- ☐ C PUT
- ☐ D DELETE



3 What does an idempotent operation ensure?

- ☐ A It is always fast
- ☐ B It can be cached
- ☐ C Multiple identical requests have the same effect as one
- ☐ D It is always safe



4 Which status code indicates that the request was successful and a resource was created?

- ☐ A 200 OK
- ☐ B 201 Created
- ☐ C 204 No Content
- ☐ D 400 Bad Request



5 Which component of an HTTP request is used to indicate the expected response media type?

- ☐ A Content-Type
- ☐ B Accept
- ☐ C Authorization
- ☐ D User-Agent



6 Which approach puts the version number directly in the URI path?

- ☐ A Header-Based Versioning
- ☐ B Query Parameter Versioning
- ☐ C URI Versioning
- ☐ D Media Type Versioning



CHECK YOURSELF



Review the concepts and examples in this module.



Think about how each concept applies to real APIs.



Revisit any topics that were challenging.



Aim for a strong understanding before moving to the lab.



Tip

Good API design leads to scalable, maintainable, and developer-friendly enterprise solutions.



Remember:

A well-designed REST API is consistent, predictable, and easy to use.

Follow the principles, use the right HTTP methods, and design with your consumers in mind.

MODULE 37 COMPLETE

PostgreSQL Database Fundamentals

You can now design normalized PostgreSQL schemas, enforce integrity with keys and constraints, index for performance, and connect from psql and Java.